

**Kubernetes
the Deltatre way**

18 MAY, 5:30 PM

The basics - Introduction to Containers and Orchestrators

RAUNO DE PASQUALE (NEWESIS)
SUPPORTED BY
CRISTIANO DEGIORGIS (DELTATRE)



newesis
e Professional | Have Fun!



Deltatre Innovation Lab



- ▶ LinkedIn:
<https://www.linkedin.com/in/rauno-de-pasquale-b075773>
- ▶ Twitter: @RaunoDepa



- ▶ LinkedIn:
<https://www.linkedin.com/in/cristianodegiorgis/>
- ▶ StackOverflow:
<https://stackoverflow.com/users/539684/crixo>

ABOUT US

- ▶ Rauno De Pasquale, Co-Founder and CTO at Newesis Srl, constantly trying to reconcile his degree in Philosophy with a passion for computer science. After almost 18 year at Deltatre, at the beginning of 2019 he creates Newesis, with the aim of simplifying the use of the most advanced services of Cloud platforms even in fields other than sports.
- ▶ Cristiano Degiorgis, An enthusiastic *lehrling* in the IT world still feeling like Alice in wonderland after so many years being around.

AGENDA

- ▶ Knowing the context and the concepts behind the use of containers is essential to be able to proceed on the road that will lead you to master the Kubernetes and Cloud Native applications.
- ▶ This initial session covers basic skills to answer questions such as:
 - ▶ what is a container image?
 - ▶ Why did anyone feel the need for an orchestrator?
 - ▶ Are there alternatives to Docker and Kubernetes?
 - ▶ How does working with containers and Kubernetes connect to traditional virtualization?
- ▶ This session has the scope of providing the basic skills to be able to orientate in subsequent sessions where the ways of creating and running applications in the Kubernetes environment will be addressed.
- ▶ Speaker: Rauno De Pasquale (Newesis) supported by Cristiano DeGiorgis (Deltatre)
- ▶ Organised by: #DeltatreLab supported by #Newesis
- ▶ Powered by: #Deltatre
- ▶ Hashtags: #DeltatreK8S #Containers #Docker #Kubernetes #meetup #webinar

WHAT THIS SESSION IS NOT

- ▶ Training on what it is and how to use Docker
 - ▶ Wait for: Monday 25-May 17:30 --> 19:30 - Kubernetes the Deltatre way: Docker in Action
- ▶ Training on what it is and how to use Kubernetes
 - ▶ Wait for:
 - ▶ Wednesday 3-Jun 17:30 --> 19:30 - Kubernetes the Deltatre way: Kubernetes basics
 - ▶ Monday 8-Jun 17:30 --> 19:30 - Kubernetes the Deltatre way: Kubernetes advanced topics & Kind
 - ▶ Monday 15-Jun 17:30 --> 19:30 - Kubernetes the Deltatre way: Kubernetes CI/CD
 - ▶ Monday 22-Jun 17:30 --> 19:30 - Kubernetes the Deltatre way: Kubernetes extensibility: CRD & Operators

INTRODUCTION TO CONTAINERS



WHAT IS A CONTAINER?

- “A container is a standard unit of software that packages up code and all its dependencies, so the application runs quickly and reliably from one computing environment to another. “ (Docker web site)
- “Containers offer a logical packaging mechanism in which applications can be abstracted from the environment in which they actually run.” (Google Cloud web site)
- “Linux containers are implementations of operating system-level virtualization for the Linux operating system.” “OS-level virtualization refers to an operating system paradigm in which the kernel allows the existence of multiple isolated user space instances. Such instances, called containers (Solaris, Docker), Zones (Solaris), virtual private servers (OpenVZ), partitions, virtual environments (VEs), virtual kernel (DragonFly BSD), or jails (FreeBSD jail or chroot jail),[1] may look like real computers from the point of view of programs running in them. “ (Wikipedia on Linux Containers and OS-Level virtualisation)
- “Isolated area of an OS with resource limits usage applied” (Nigel Poulton, book “Docker Deep Dive”)

A CONTAINER IS A GROUP OF PROCESSES ...

Containers

Processes

host:aks-default-37063012-3

Hide Controls Showing 1 – 18 of 18 containers

NAME

mysql_wordpress-mysql-66594

mysql_onlinelandpg-mysql-8bc

prestashop_prestashop-fc8d59bc7-vx2mr

planetps_planetps-dfd577f4c-4

agent_datadog-2ptms

webform_webform-5b9bbcbcf4

process-agent_datadog-2ptms

phpmyadmin_phpmyadmin-f66

onlinelandpg_onlinelandpg-55a

coredns_coredns-6c66fc4fcb-b7

nginx_proxy-zeca-zecaproducti

nginx_proxy-donazioni-7cdc788

nginx_proxy-6756d9d8d5-kn6q

nginx_proxy-miolascito-5d5dbc

nginx_proxy-disabilita-57c5465

autoscaler_coredns-autoscaler

kube-proxy_kube-proxy-bds4s

nginx_planetps-dfd577f4c-4qlg

prestashop_prestashop-fc8d59bc7-vx2mr

aks-default-37063012-313 days ago

newesisprdregistry.azurecr.io/psauto/psauto

container_id:153e9b74c213pod_name:prestashop-fc8d59bc7-vx2mrhost:aks-default-37063012-3aksengineversion:v0.43.0-aksavailability-zone:northeurope+27

PROCESS LIST

COMMAND	PID	PPID	CPU %	RSS MEMORY
◀ apache2 -DFOREGROUND	86799	86781	0 %	63 MB
▶ apache2 -DFOREGROUND	87686	86799	0 %	103 MB
▶ apache2 -DFOREGROUND	85116	86799	0 %	96 MB
▶ apache2 -DFOREGROUND	103759	86799	0 %	96 MB

Metrics

Logs

Traces

Network

12hPast 12 Hours

Total CPU (%)6e-3 % (0 %)

RSS Memory134.1 MiB (134.1 MiB)

0.25

0.2

0.15

0.1

0.05

0

09:0012:0015:0018:00

160

128

96

64

32

0

09:0012:0015:0018:00

Containers

Processes

host:aks-default-37063012-3

Hide Controls Showing 1 – 50 of 50 processes

Scatter Plot

Timeseries

Graph

Total CPU (%)

Avg

RSS Memory

log

2G

0.5G

0.12G

0.031G

0.0078G

0.002G

0.00049G

0.00012G

0.000031G

0.0000076G

0.0000019G

1e-3

apache2

aks-default-37063012-3root13 days ago

86799

pod_name:prestashop-fc8d59bc7-vx2mrhost:aks-default-37063012-3aksengineversion:v0.43.0-aksavailability-zone:northeuropecloud_provider:azure+28

FULL COMMAND

apache2 -DFOREGROUND

12hPast 12 Hours

User CPU (%)490077 n% (0 %)

RSS Memory59.9 MiB (59.9 MiB)

0.01

8e-3

6e-3

4e-3

2e-3

0

09:0012:0015:0018:00

64

48

32

16

0

09:0012:0015:0018:00

PROCESS

mysqld

mysqld

kubelet --enable-server --node-label

systemd-journald

agent start

dotnet webform.dll

dockerd -H fd:// --storage-driver=ov

apache2 -DFOREGROUND

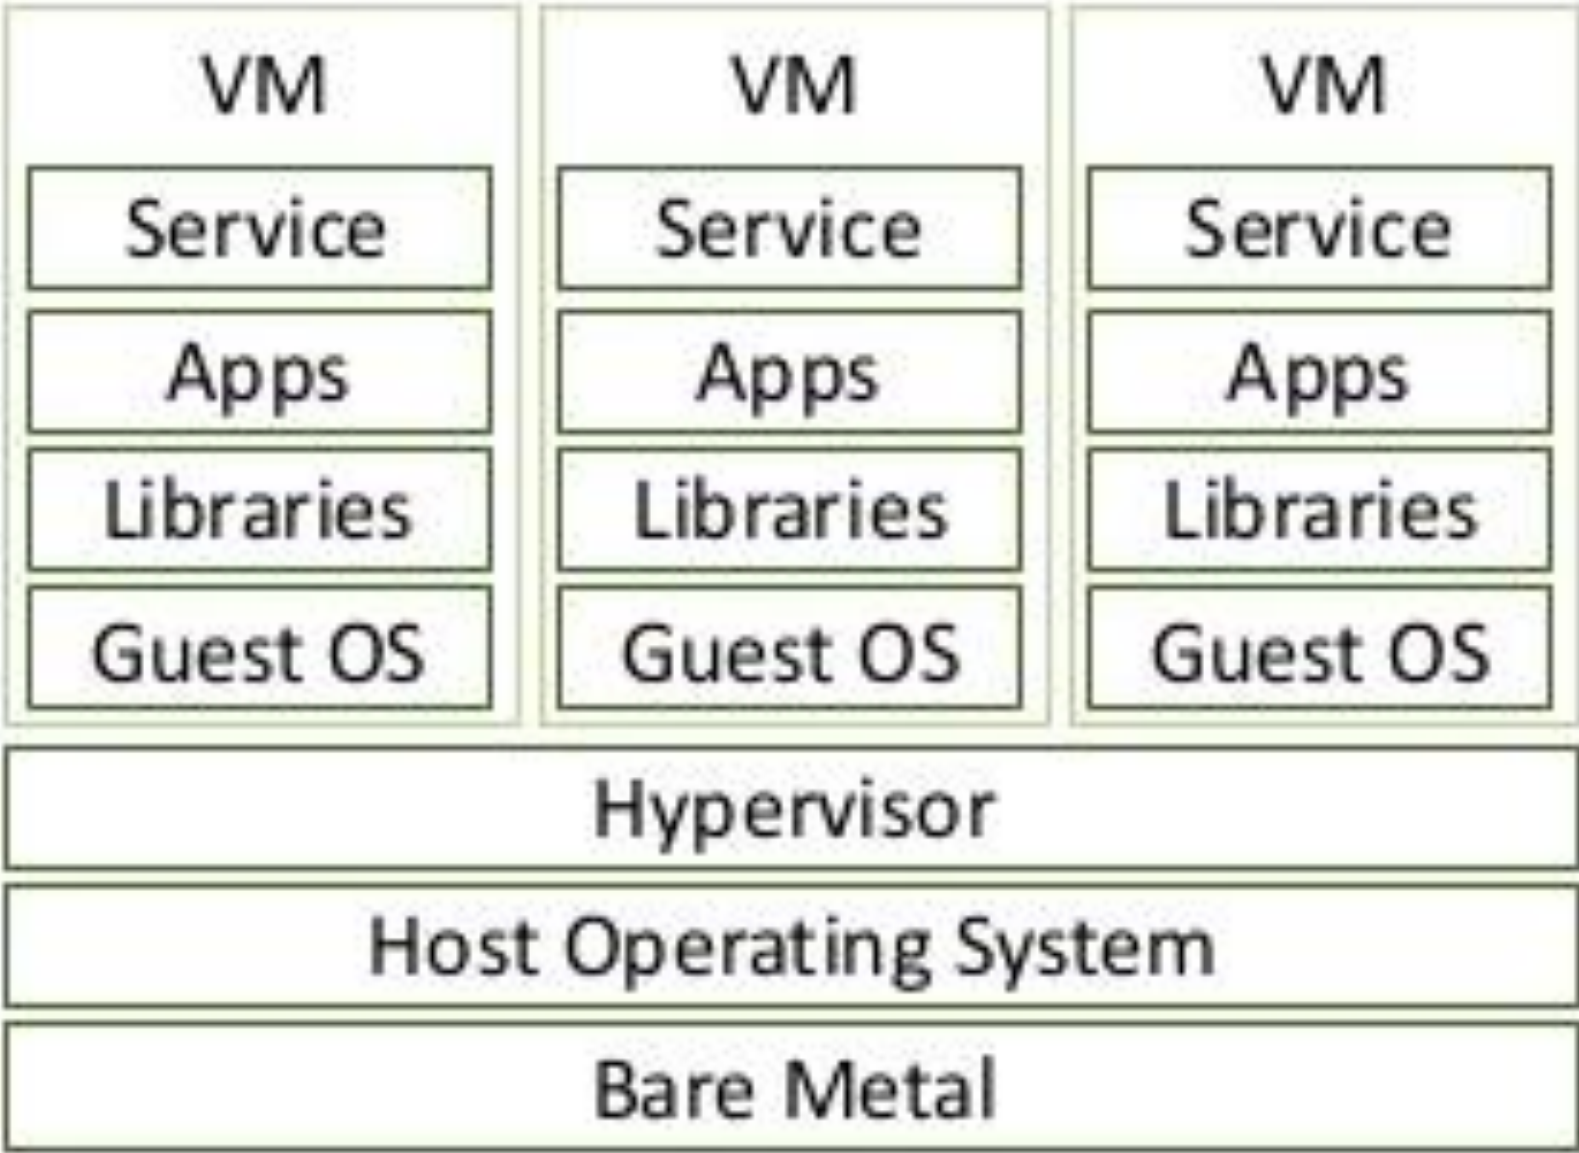
apache2 -DFOREGROUND

... RESTRICTED TO A PRIVATE NAMESPACE

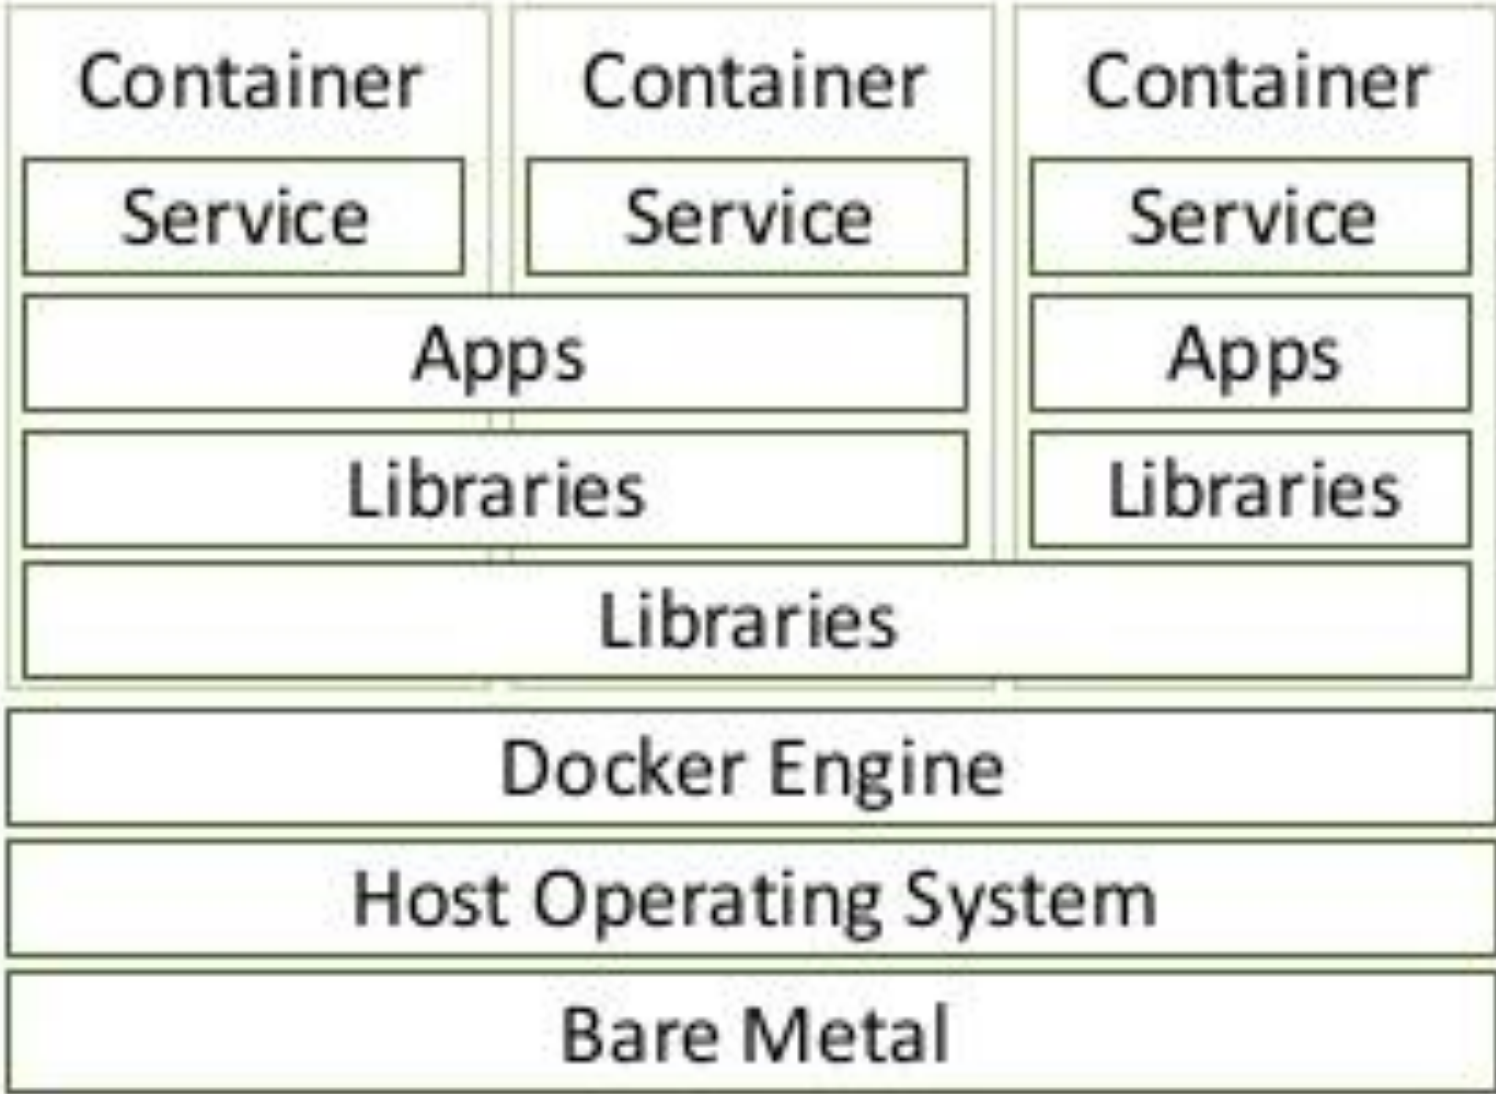
- “Namespaces are a feature of the Linux kernel that partitions kernel resources such that one set of processes sees one set of resources while another set of processes sees a different set of resources. The feature works by having the same namespace for a set of resources and processes, but those namespaces refer to distinct resources. Resources may exist in multiple spaces. Examples of such resources are process IDs, hostnames, user IDs, file names, and some names associated with network access, and interprocess communication. “ (Wikipedia – Linux namespaces)
- “cgroups (abbreviated from control groups) is a Linux kernel feature that limits, accounts for, and isolates the resource usage (CPU, memory, disk I/O, network, etc.) of a collection of processes..” (Wikipedia - Cgroups)
- “cgroups, which stands for control groups, are a kernel mechanism for limiting and measuring the total resources used by a group of processes running on a system. For example, you can apply CPU, memory, network or IO quotas. cgroups were originally developed by Paul Menage and Rohit Seth of Google, and their first features were merged into Linux 2.6.24.” (Duncan Macrae - How Linux Kernel Cgroups And Namespaces Made Modern Containers Possible)
- “Namespaces are a kernel mechanism for limiting the visibility that a group of processes has of the rest of a system. For example you can limit visibility to certain process trees, network interfaces, user IDs or filesystem mounts. namespaces were originally developed by Eric Biederman, and the final major namespace was merged into Linux 3.8.” (Duncan Macrae - How Linux Kernel Cgroups And Namespaces Made Modern Containers Possible)

CONTAINERS VS VIRTUAL MACHINES

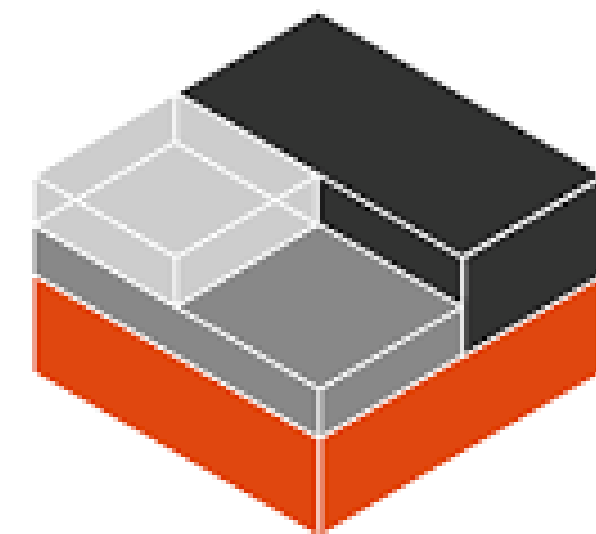
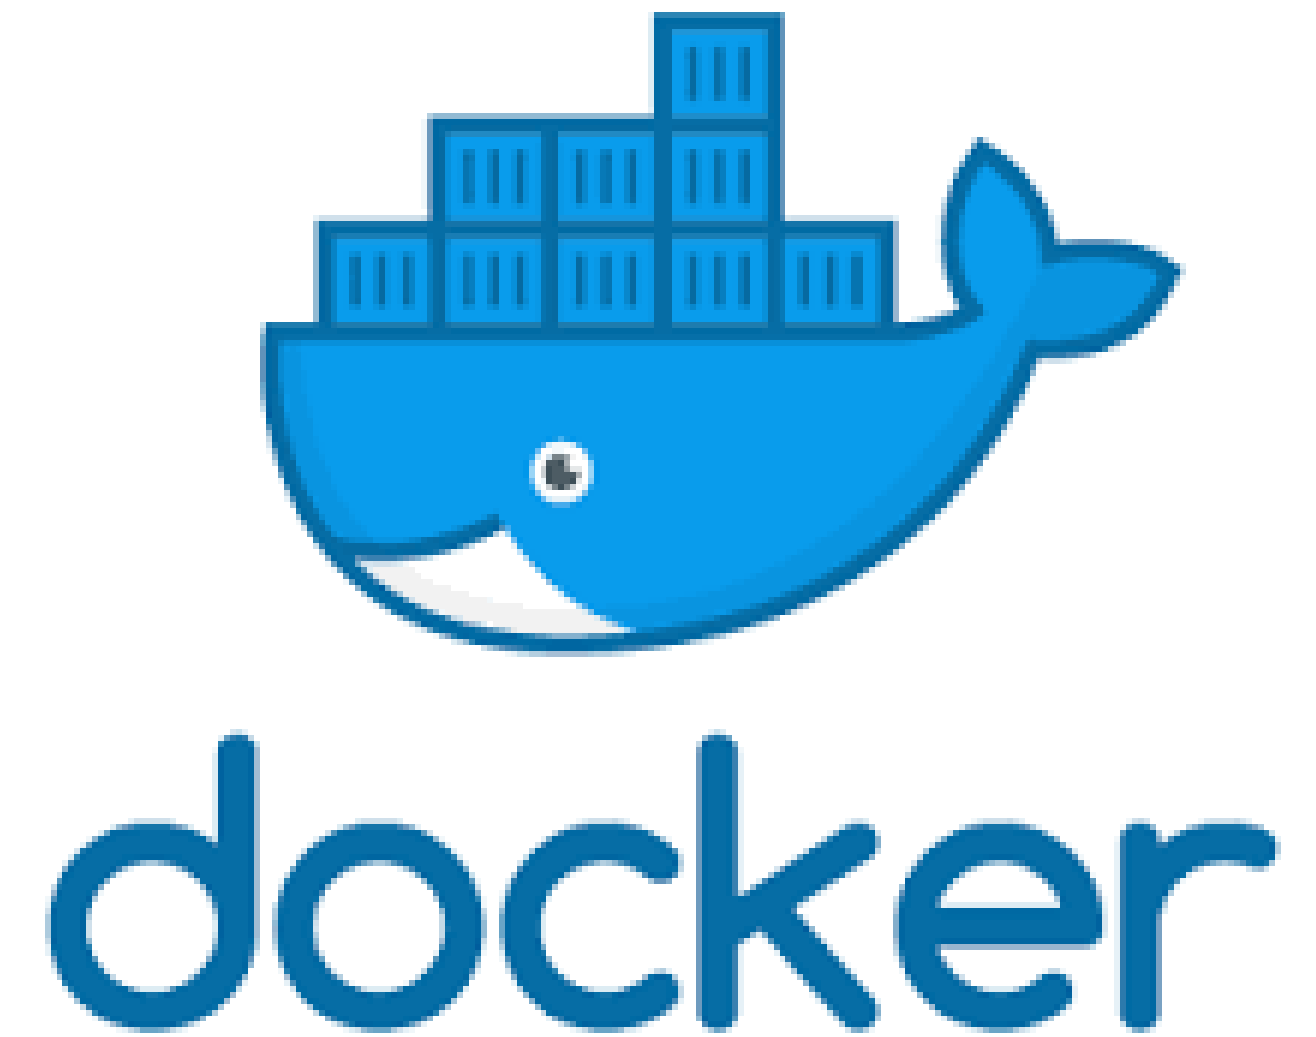
VMs



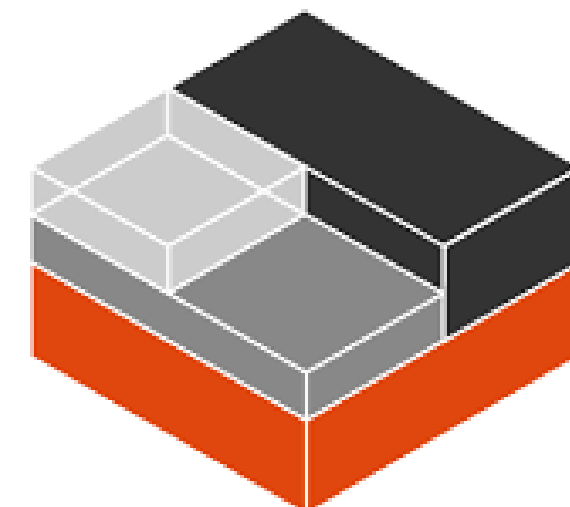
Containers



DOES CONTAINER MEAN DOCKER?



LXC



LXD



Hyper-V Container

OK CONTAINERS BUT WHY ORCHESTRATORS?

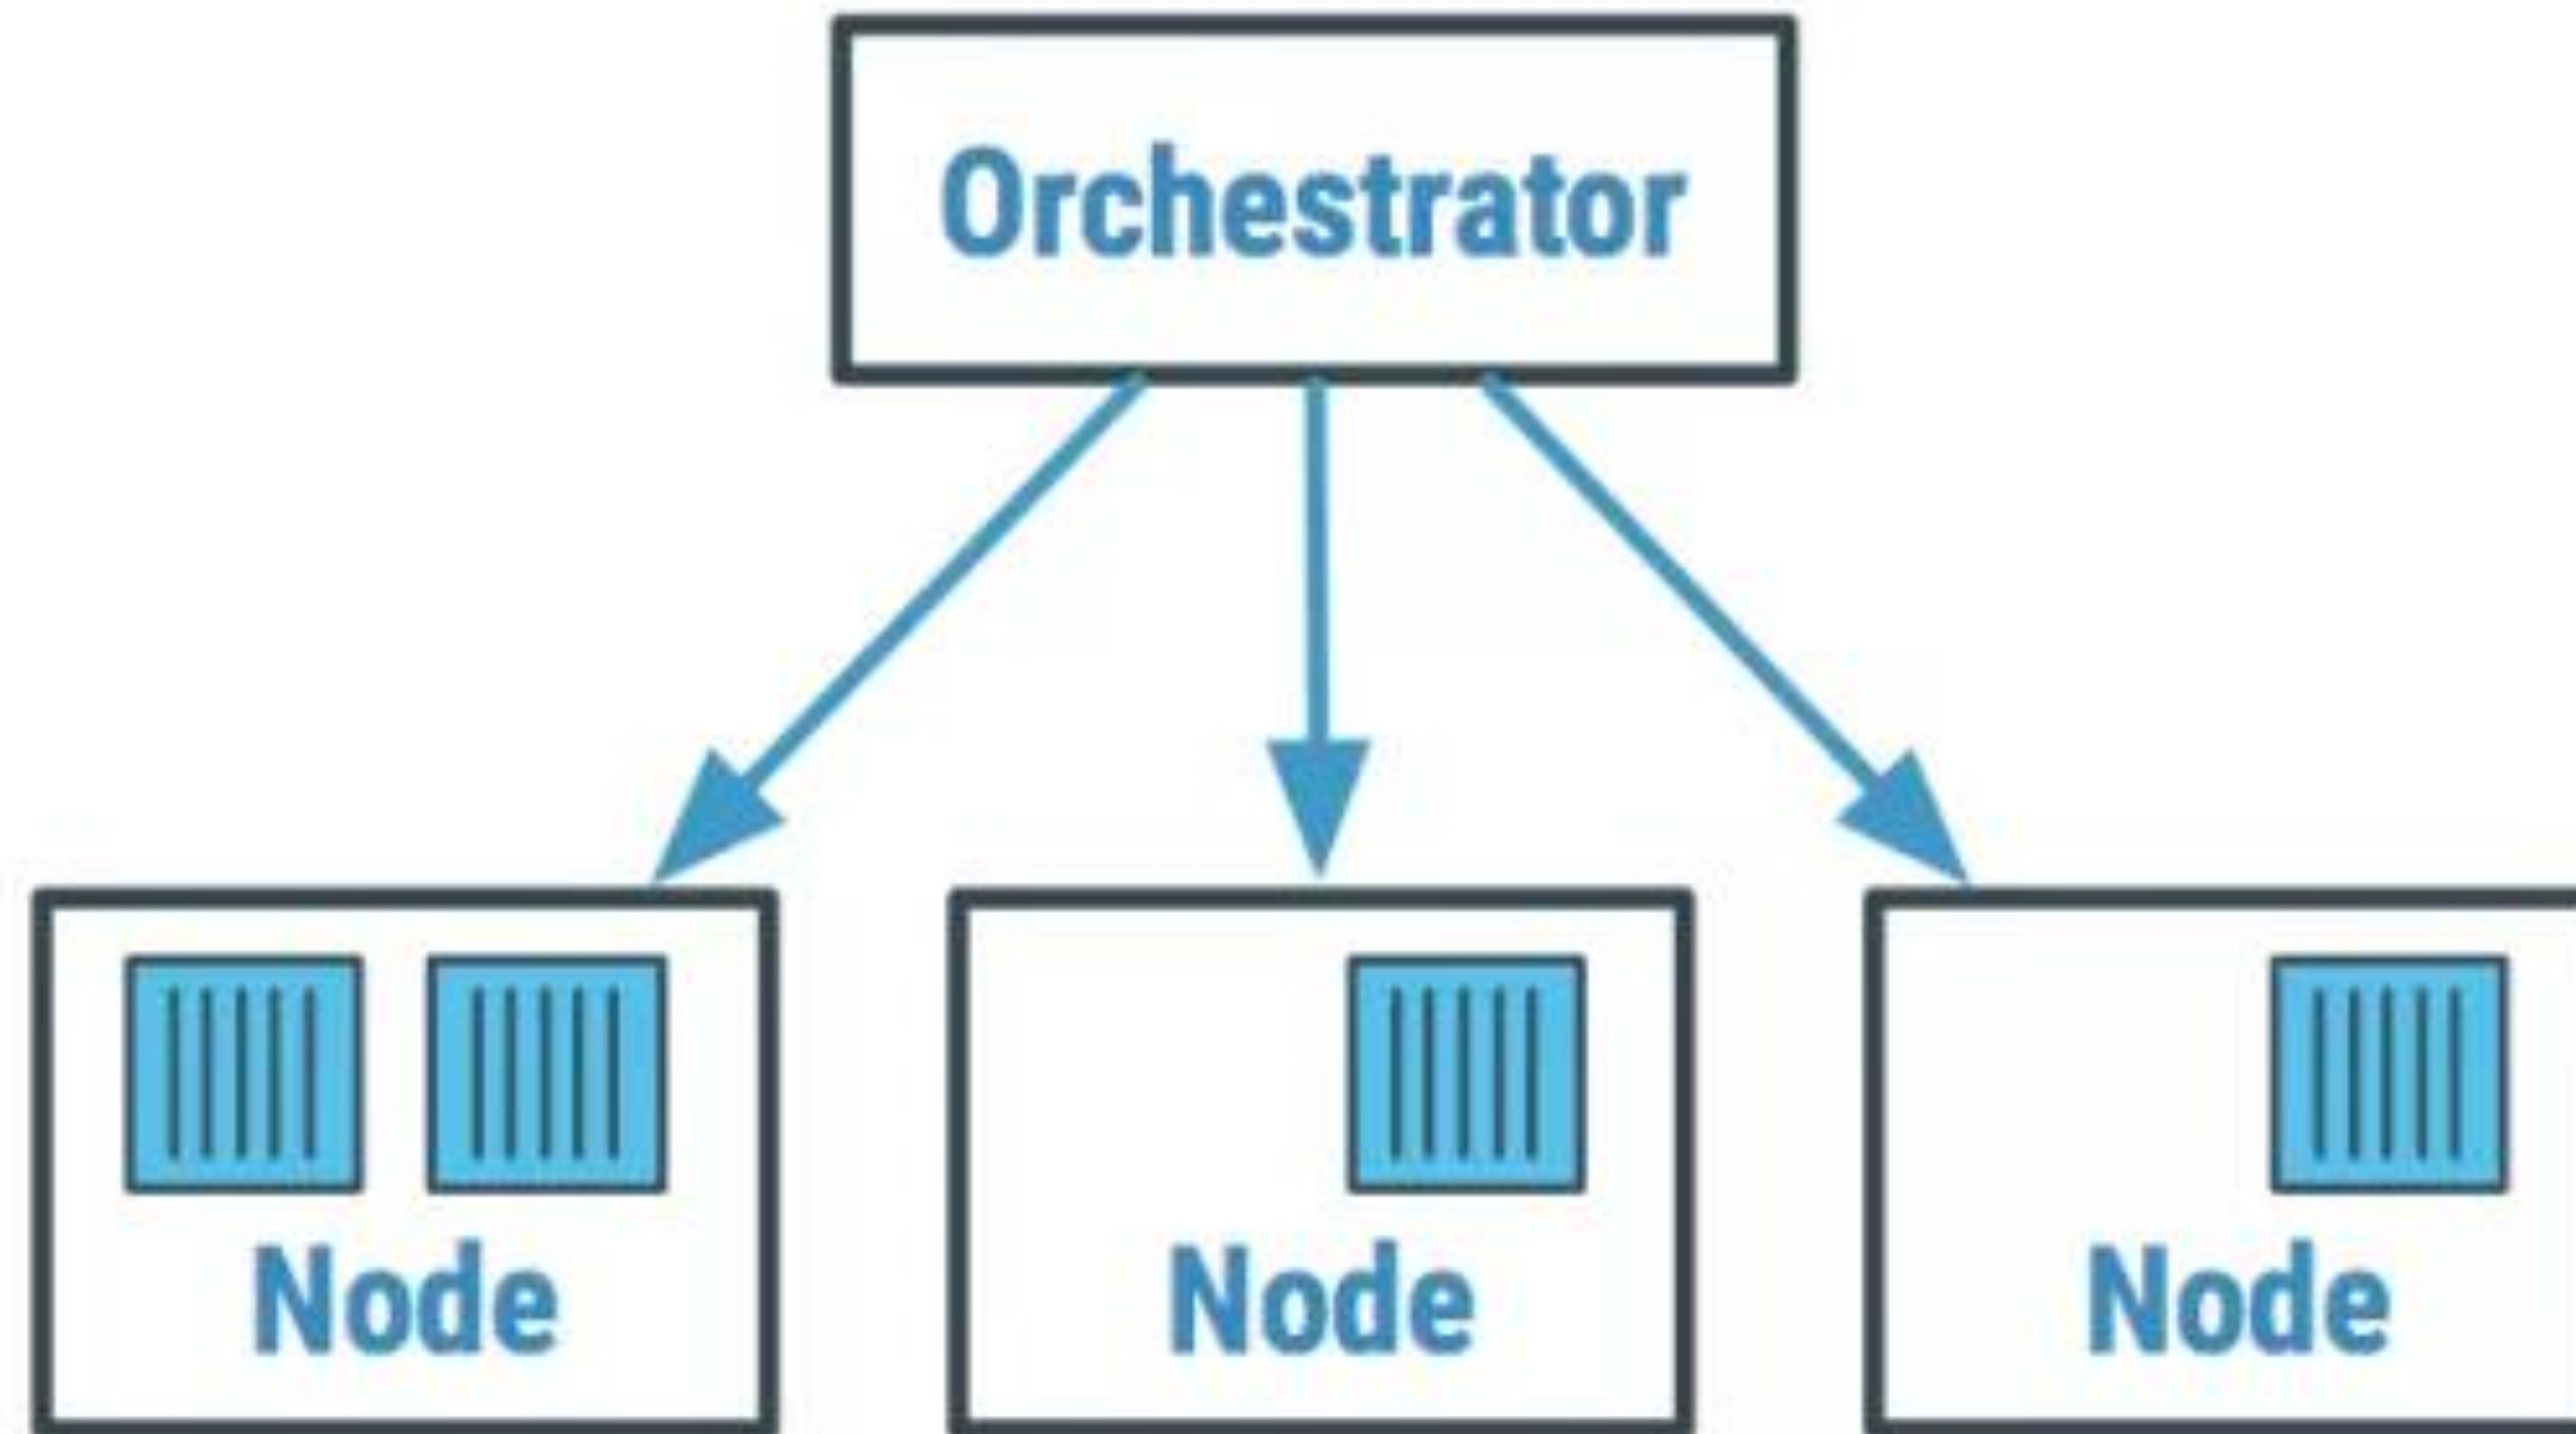


CommitStrip.com





ORCHESTRATORS



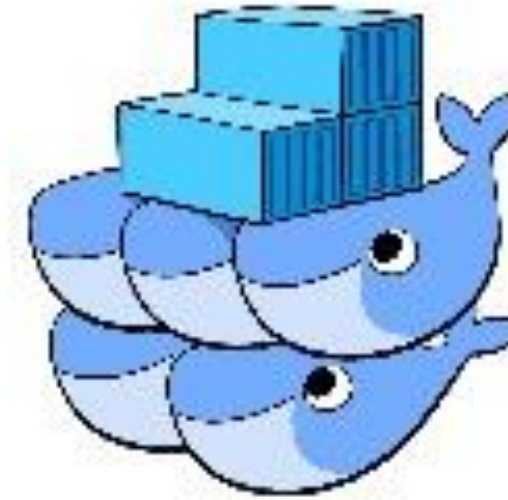
DOES ORCHESTRATOR MEAN KUBERNETES?

Container Orchestration Tools



MESOS

Marathon (Mesosphere)



Docker Swarm



Nomad (HashiCorp)



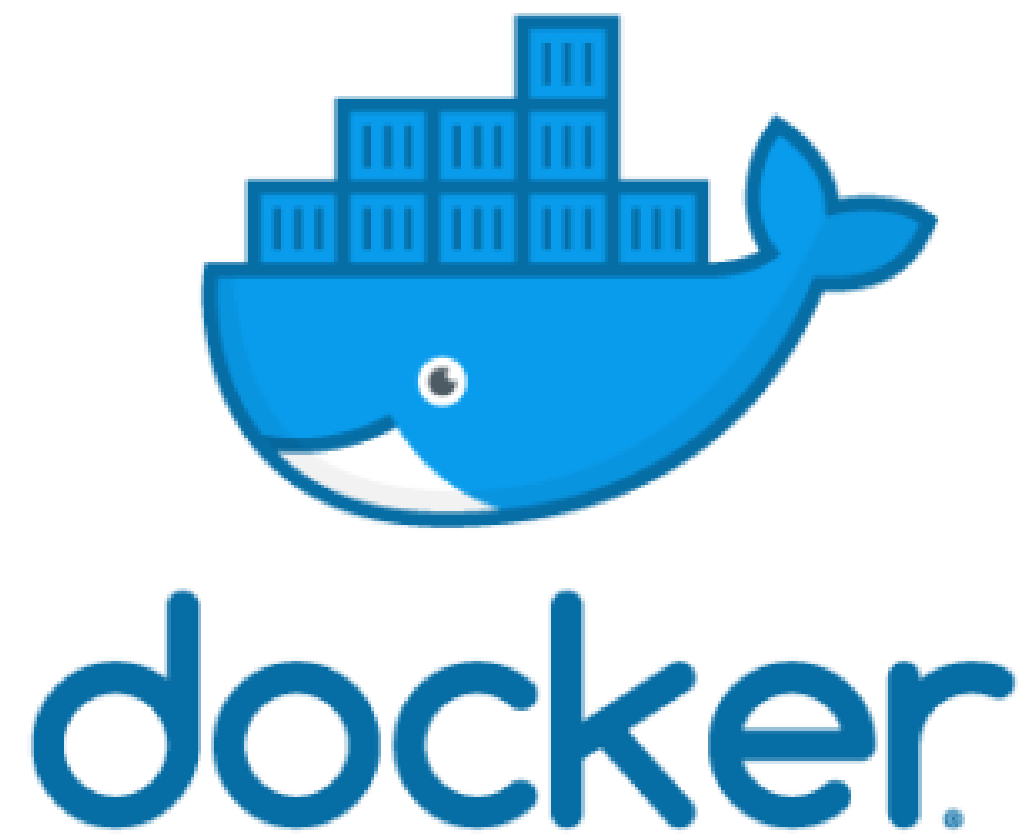
Kubernetes

WHY DOCKER AND KUBERNETES?



CommitStrip.com

DELTATRE STRATEGY



- Docker and Kubernetes have the larger communities and larger adoption
- Fully supported by all major Cloud providers
- Fully supported for an OnPremises configuration
- Part of the Open Containers Initiatives
- Part of the Cloud Native Computing Foundation
- Docker supports Kubernetes (now part of the Enterprise Edition)
- Docker support migration from Swarm to Kubernetes
- Google Borg as foundation of Kubernetes



CLOUD AGNOSTIC

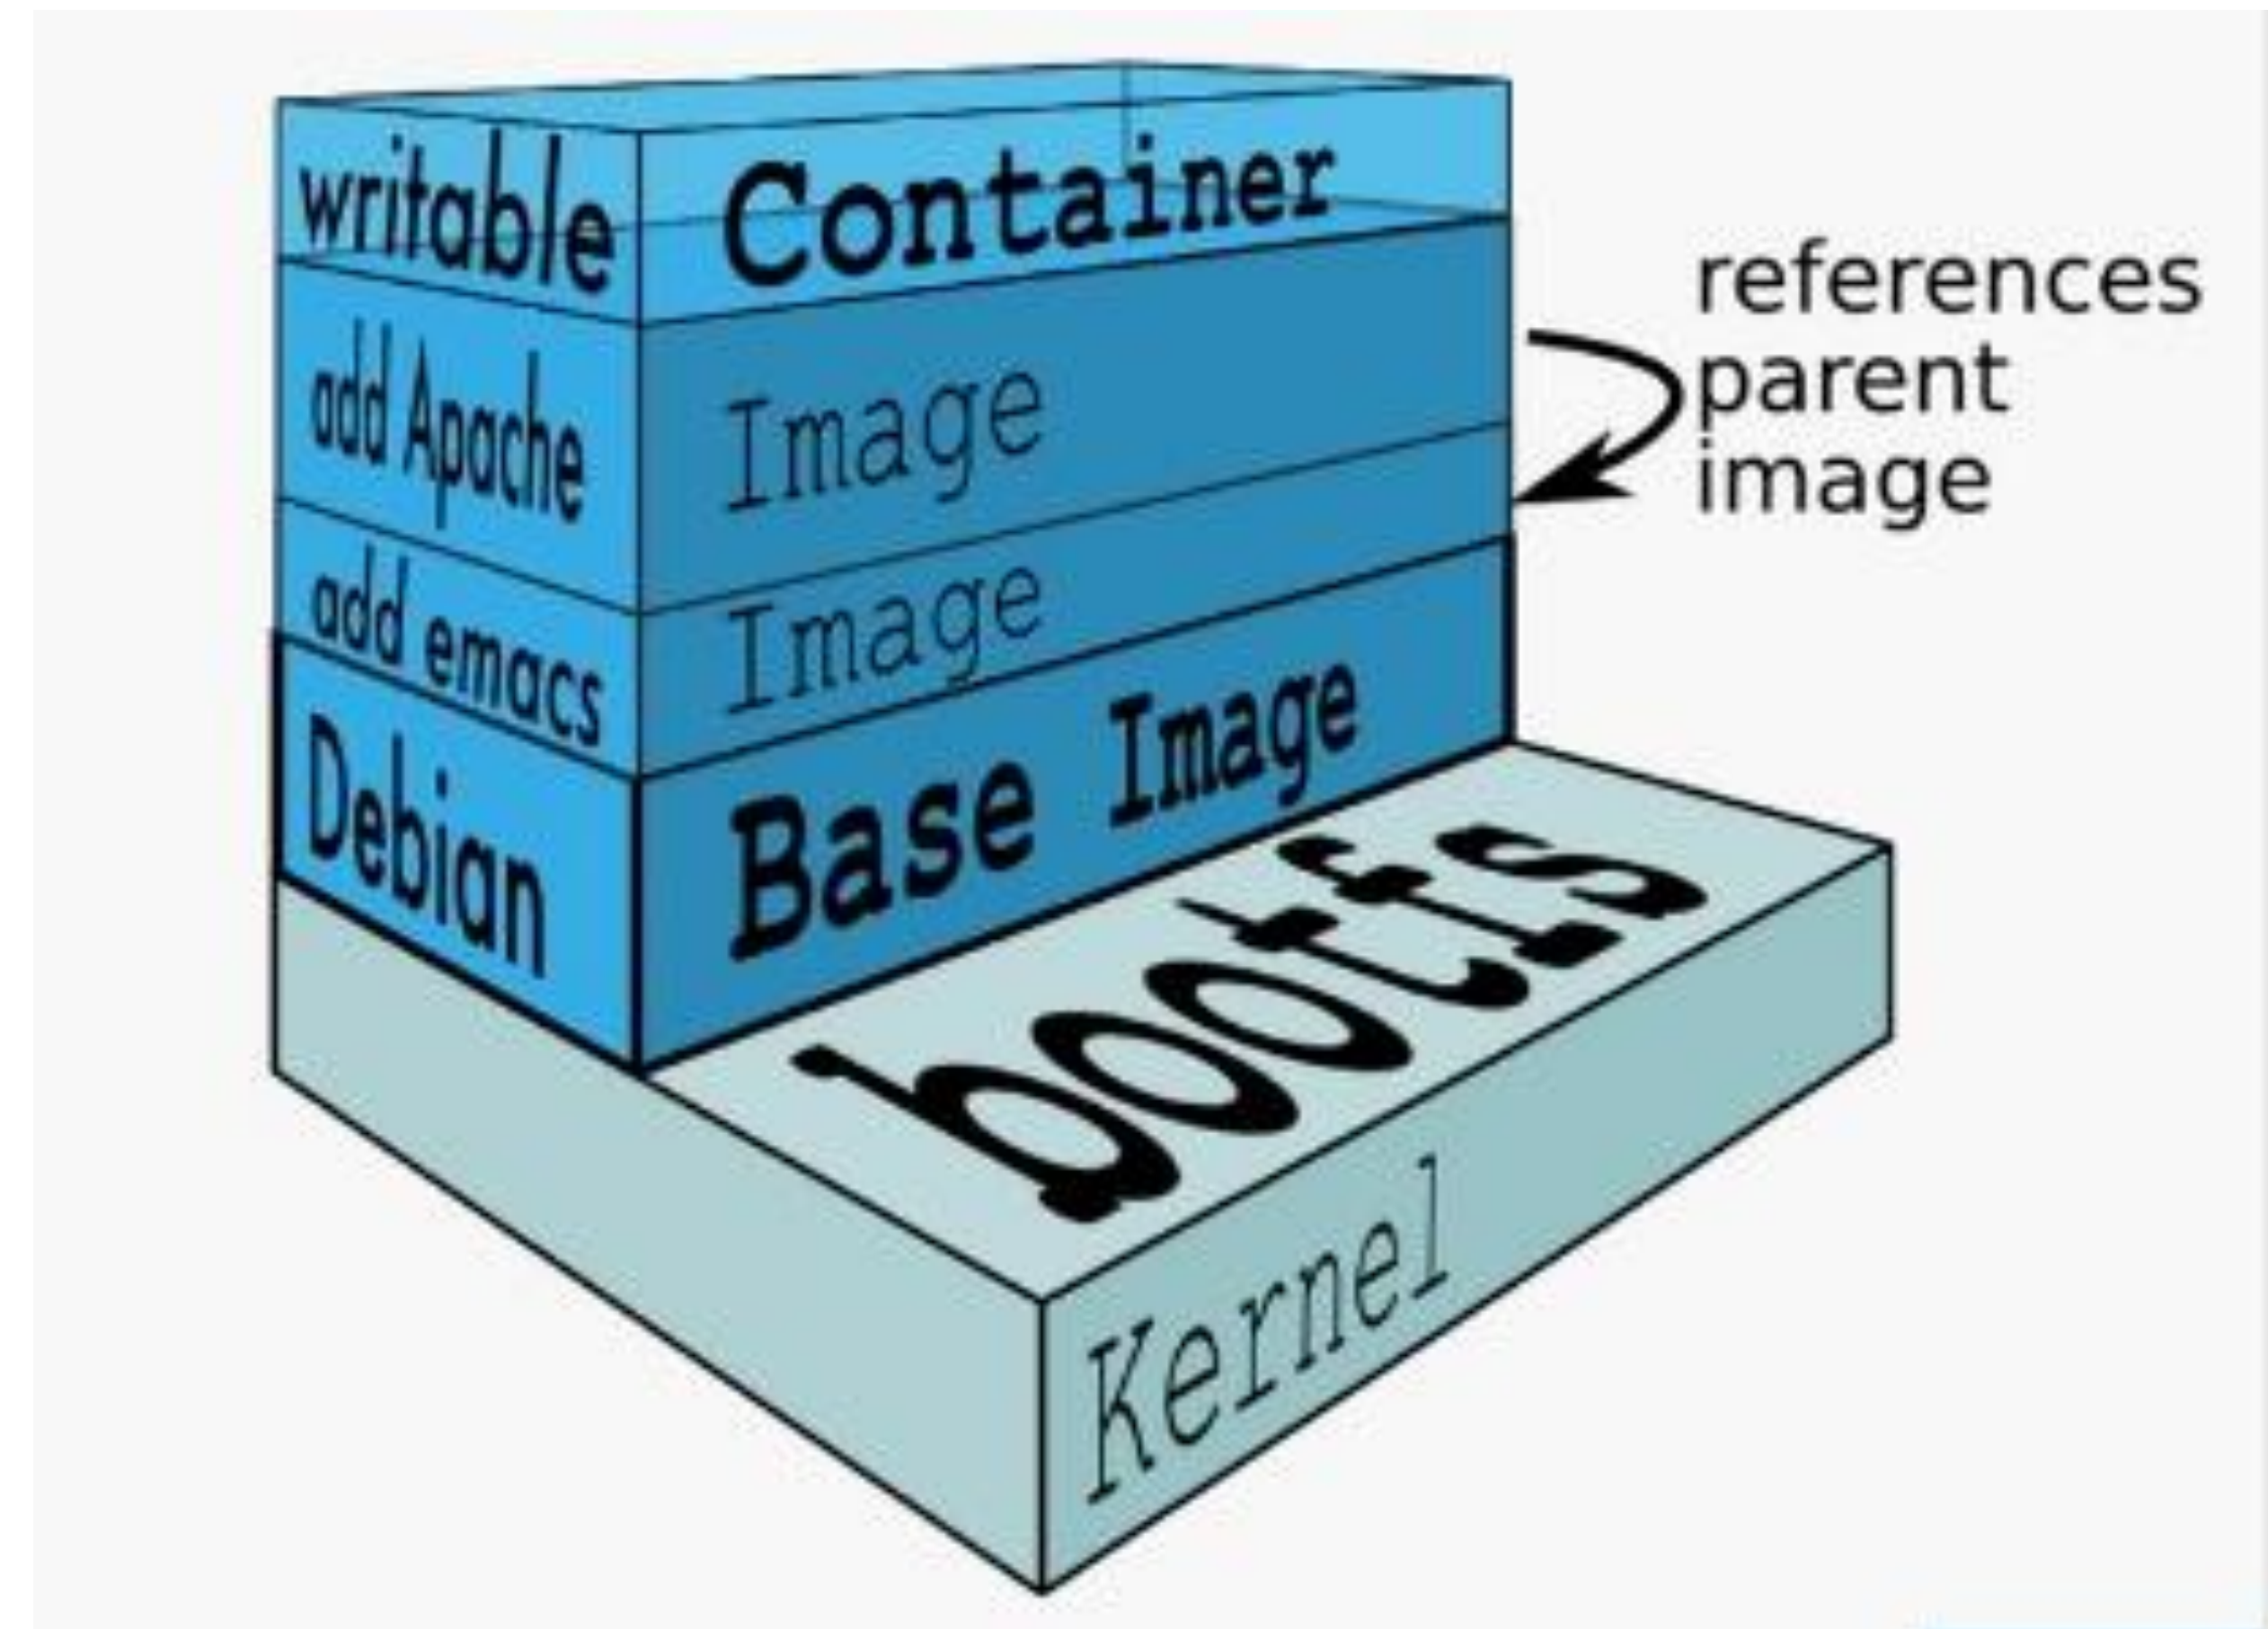
Portable Solutions

- Reusable components and products must be Cloud Agnostic
- Container images able to run on Linux OS
 - NodeJS
 - .Net Core
- Docker images and Kubernetes based deployments
- MongoDB and in general intensive IO applications to be installed into VMs and not as containers
- Usage of PaaS only if replaceable with alternatives (e.g. CosmosDB in Azure is ok if development is done to preserve compatibility with MongoDB)

Docker

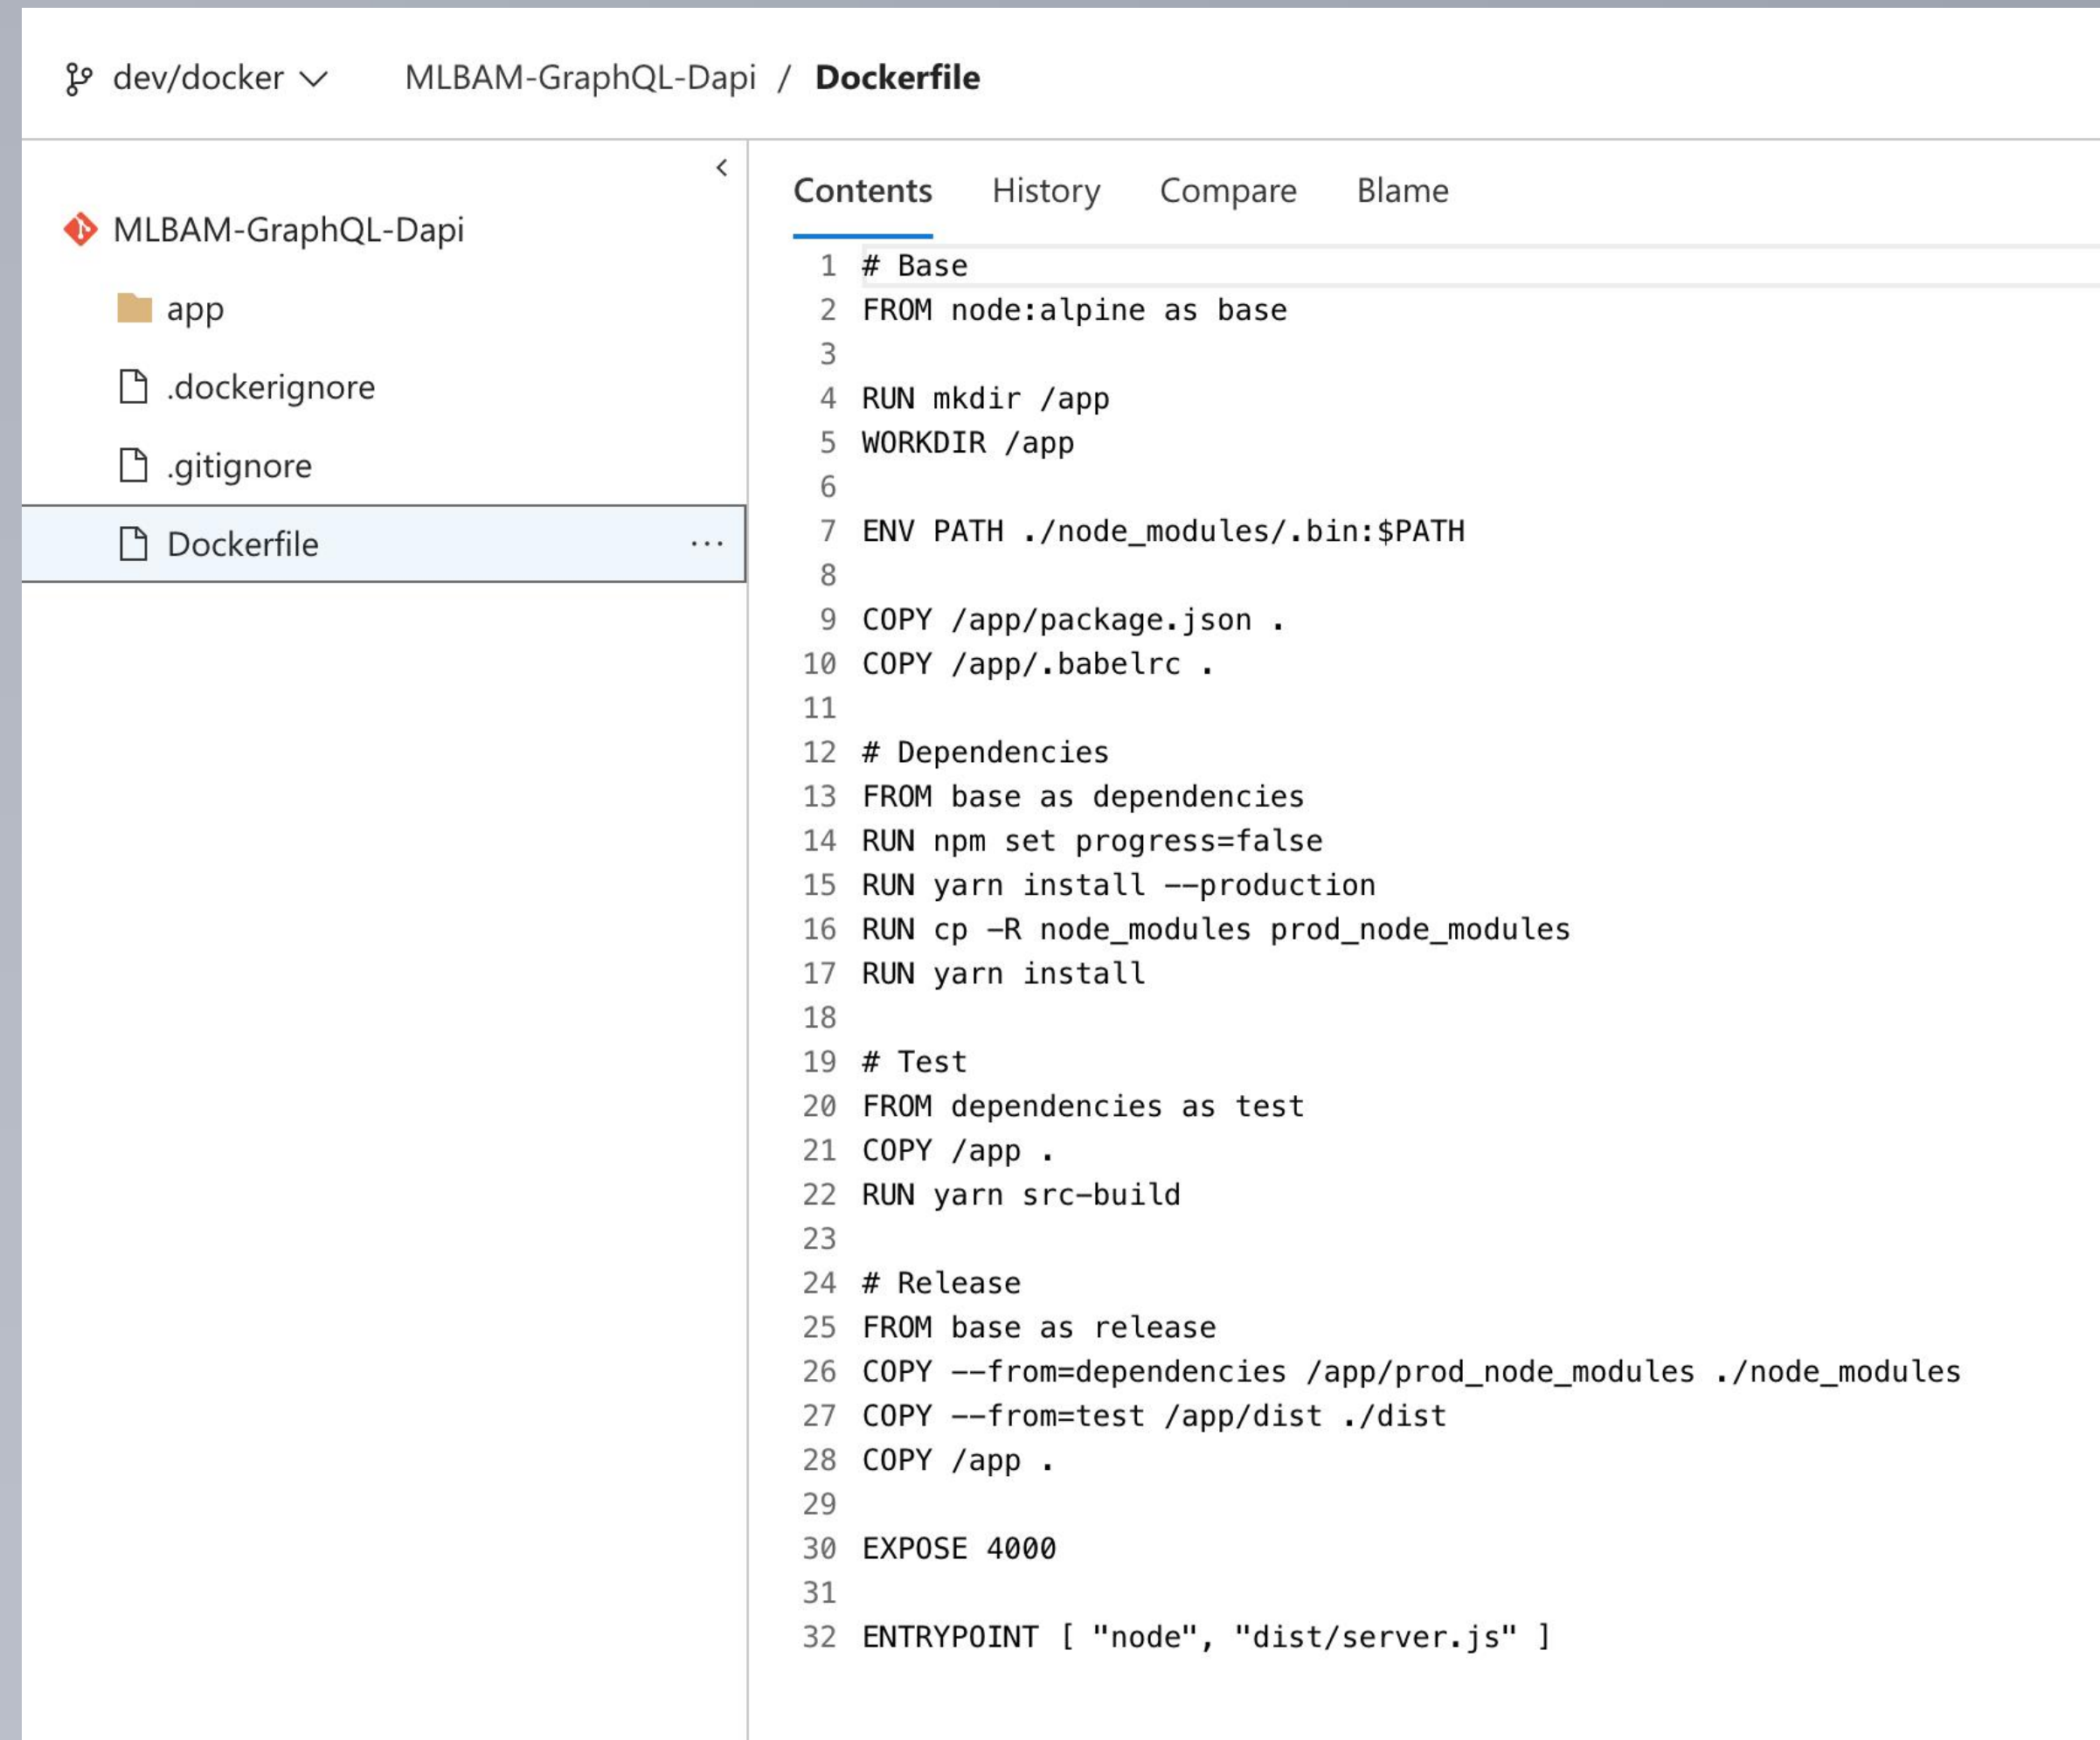
DOCKER BASICS

- Dockerfile
 - Source code of an image
- Image
 - Immutable package of application and its dependencies
 - Composed by multiple layers
- Container
 - Running instance of an image
- Registry
 - Repository of images
- Docker Daemon
 - Build images
 - Run Containers
- Docker CLI



DOCKERFILE

Image build instruction



The screenshot shows a code editor interface with a sidebar on the left and a main editor area on the right. The sidebar displays the file structure of a project named 'MLBAM-GraphQL-Dapi', including an 'app' folder and files like '.dockerignore', '.gitignore', and 'Dockerfile'. The 'Dockerfile' is selected and highlighted. The main editor area shows the contents of the Dockerfile, which is a multi-stage build. It starts with a 'Base' stage using 'node:alpine' as the base image, followed by setting up the environment, copying files, and installing dependencies. The 'Dependencies' stage uses 'base' as the parent image to install npm dependencies. The 'Test' stage uses 'dependencies' as the parent image to run a build. The 'Release' stage uses 'base' as the parent image to copy the production dependencies and the built application. Finally, the 'ENTRYPOINT' is set to run 'node' with the path to the application's server.js file.

```
1 # Base
2 FROM node:alpine as base
3
4 RUN mkdir /app
5 WORKDIR /app
6
7 ENV PATH ./node_modules/.bin:$PATH
8
9 COPY /app/package.json .
10 COPY /app/.babelrc .
11
12 # Dependencies
13 FROM base as dependencies
14 RUN npm set progress=false
15 RUN yarn install --production
16 RUN cp -R node_modules prod_node_modules
17 RUN yarn install
18
19 # Test
20 FROM dependencies as test
21 COPY /app .
22 RUN yarn src-build
23
24 # Release
25 FROM base as release
26 COPY --from=dependencies /app/prod_node_modules ./node_modules
27 COPY --from=test /app/dist ./dist
28 COPY /app .
29
30 EXPOSE 4000
31
32 ENTRYPOINT [ "node", "dist/server.js" ]
```

- A dockerfile contains the instruction for the docker build process on how to create a new image
- Build of an image is done by executing command inside a container
- A container is the execution of an image
- Multi-Stage builds should be used to optimise image creation process and image size

Kubernetes

KUBERNETES – THE ORIGIN



kubernetes

[——— 8 ———]

- ▶ Greek for “Helmsman”; also the root of the word “governor” and “cybernetic”
- ▶ Orchestrator for containers
- ▶ Builds on Docker containers
- ▶ Also supporting other container technologies
- ▶ Multi-cloud and bare-metal environments
- ▶ Inspired and informed by Google’s experiences and internal systems
- ▶ 100% Open Source, written in Go
- ▶ Created by three Google employees initially during the summer of 2014; grew exponentially and became the first project to get donated to the CNCF
- ▶ Release 1.0 21st July 2015

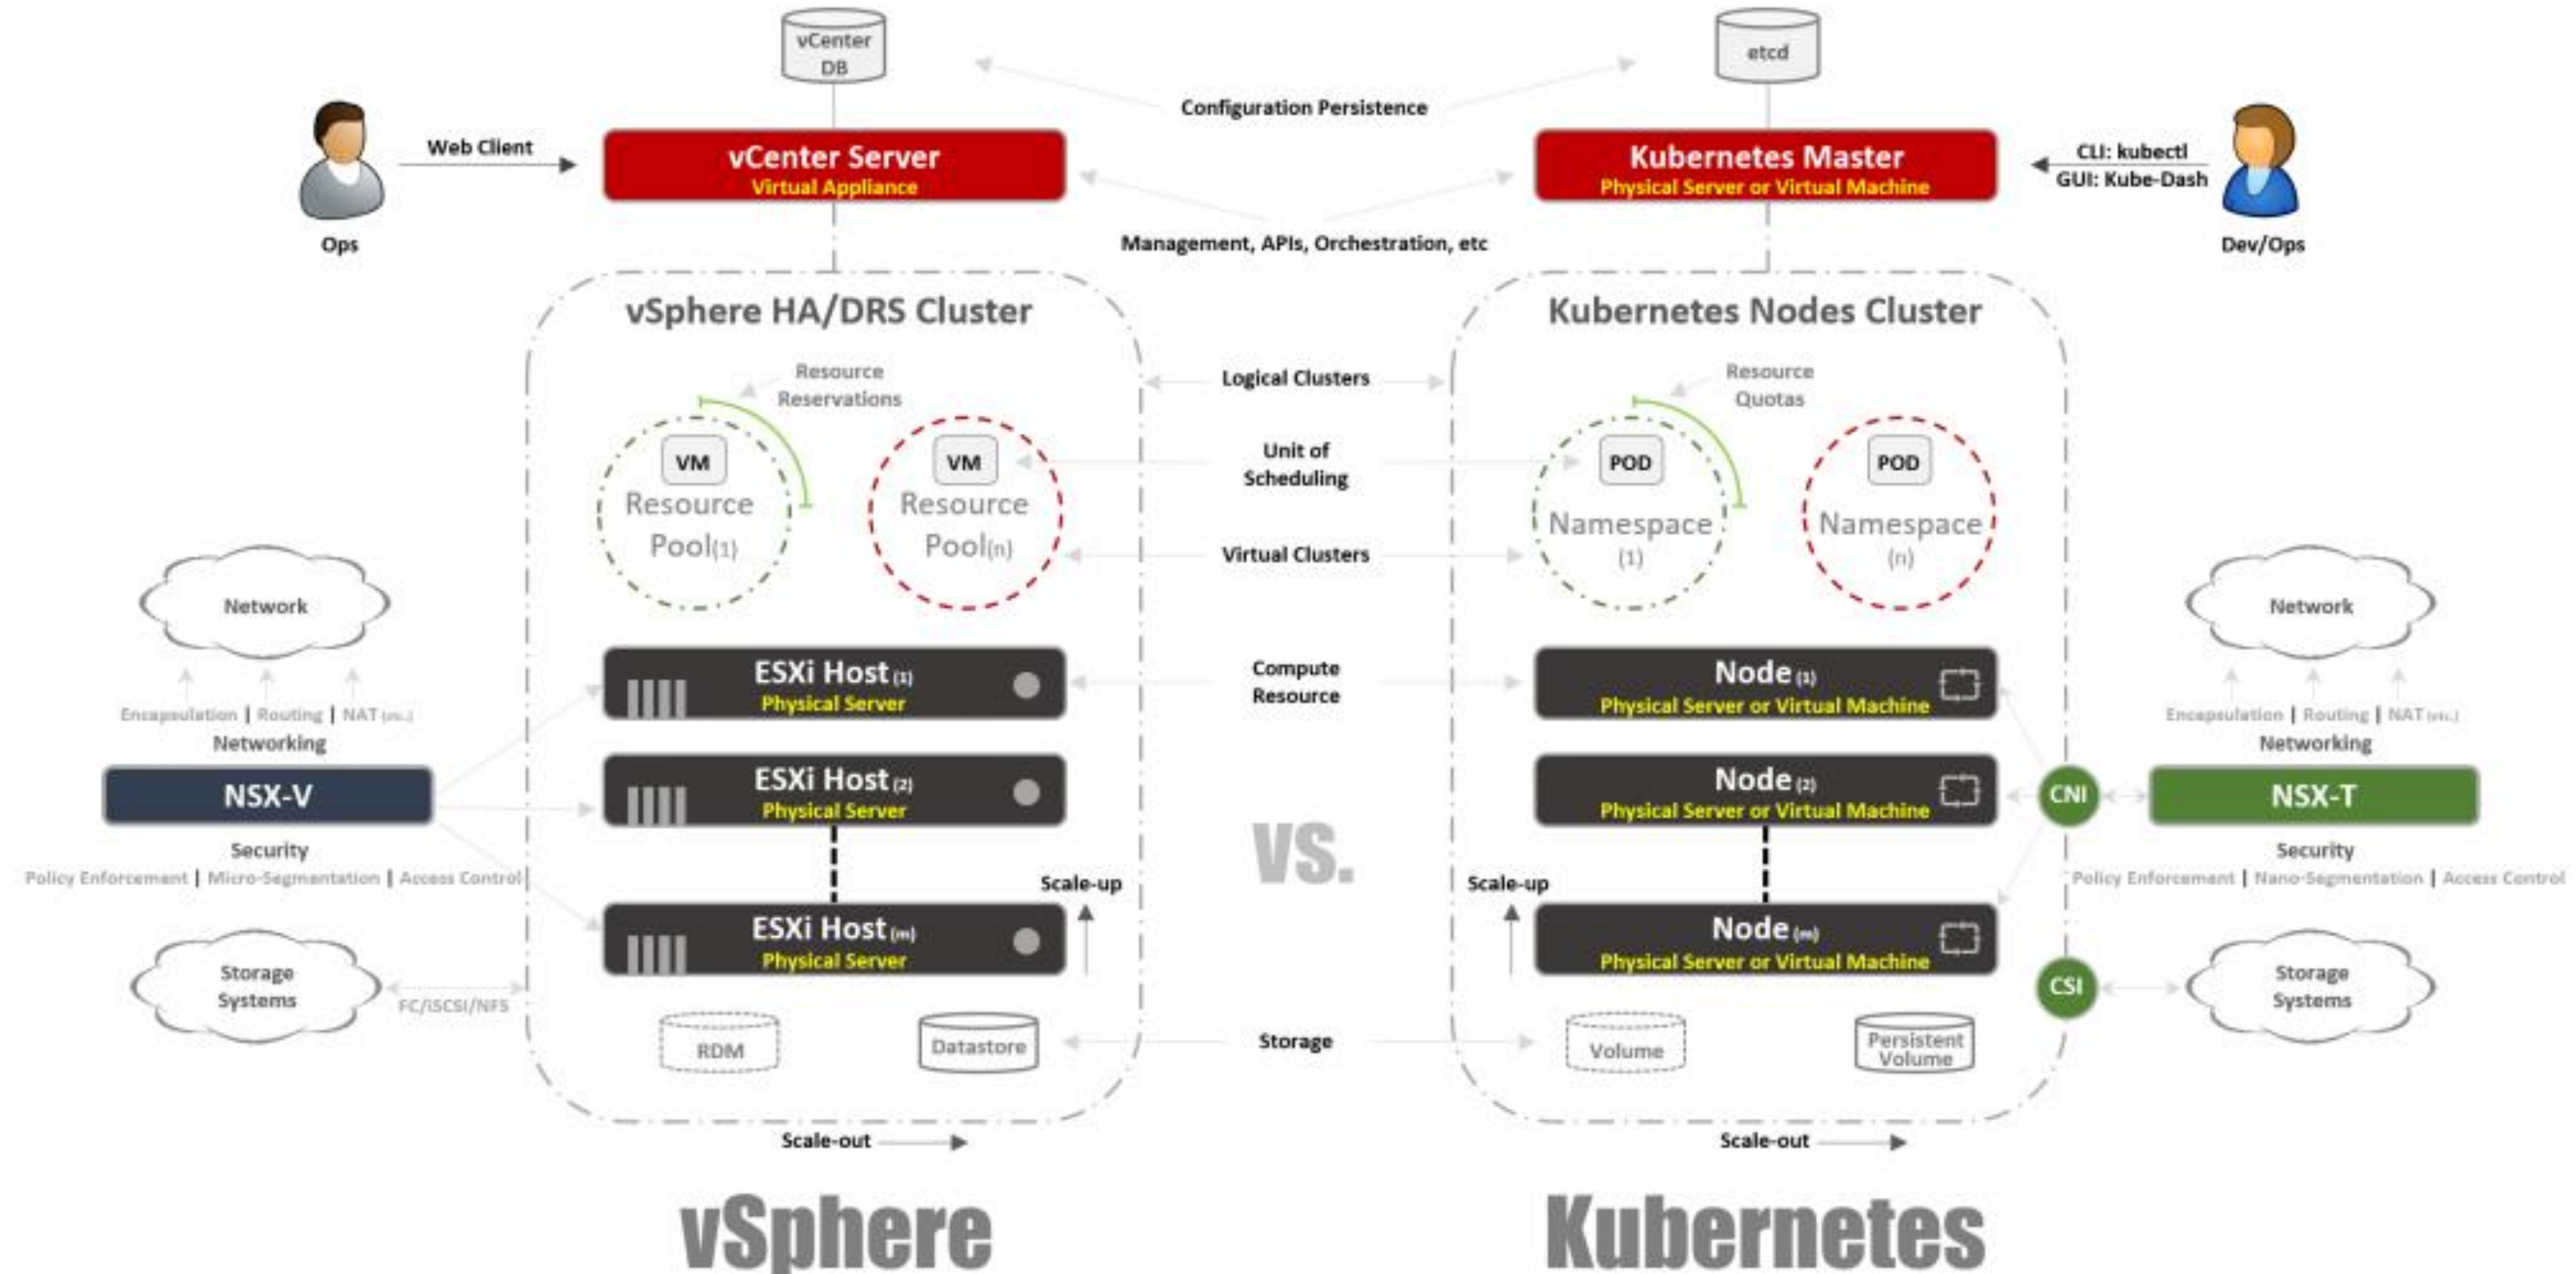
KUBERNETES – THE BASIC CONCEPTS



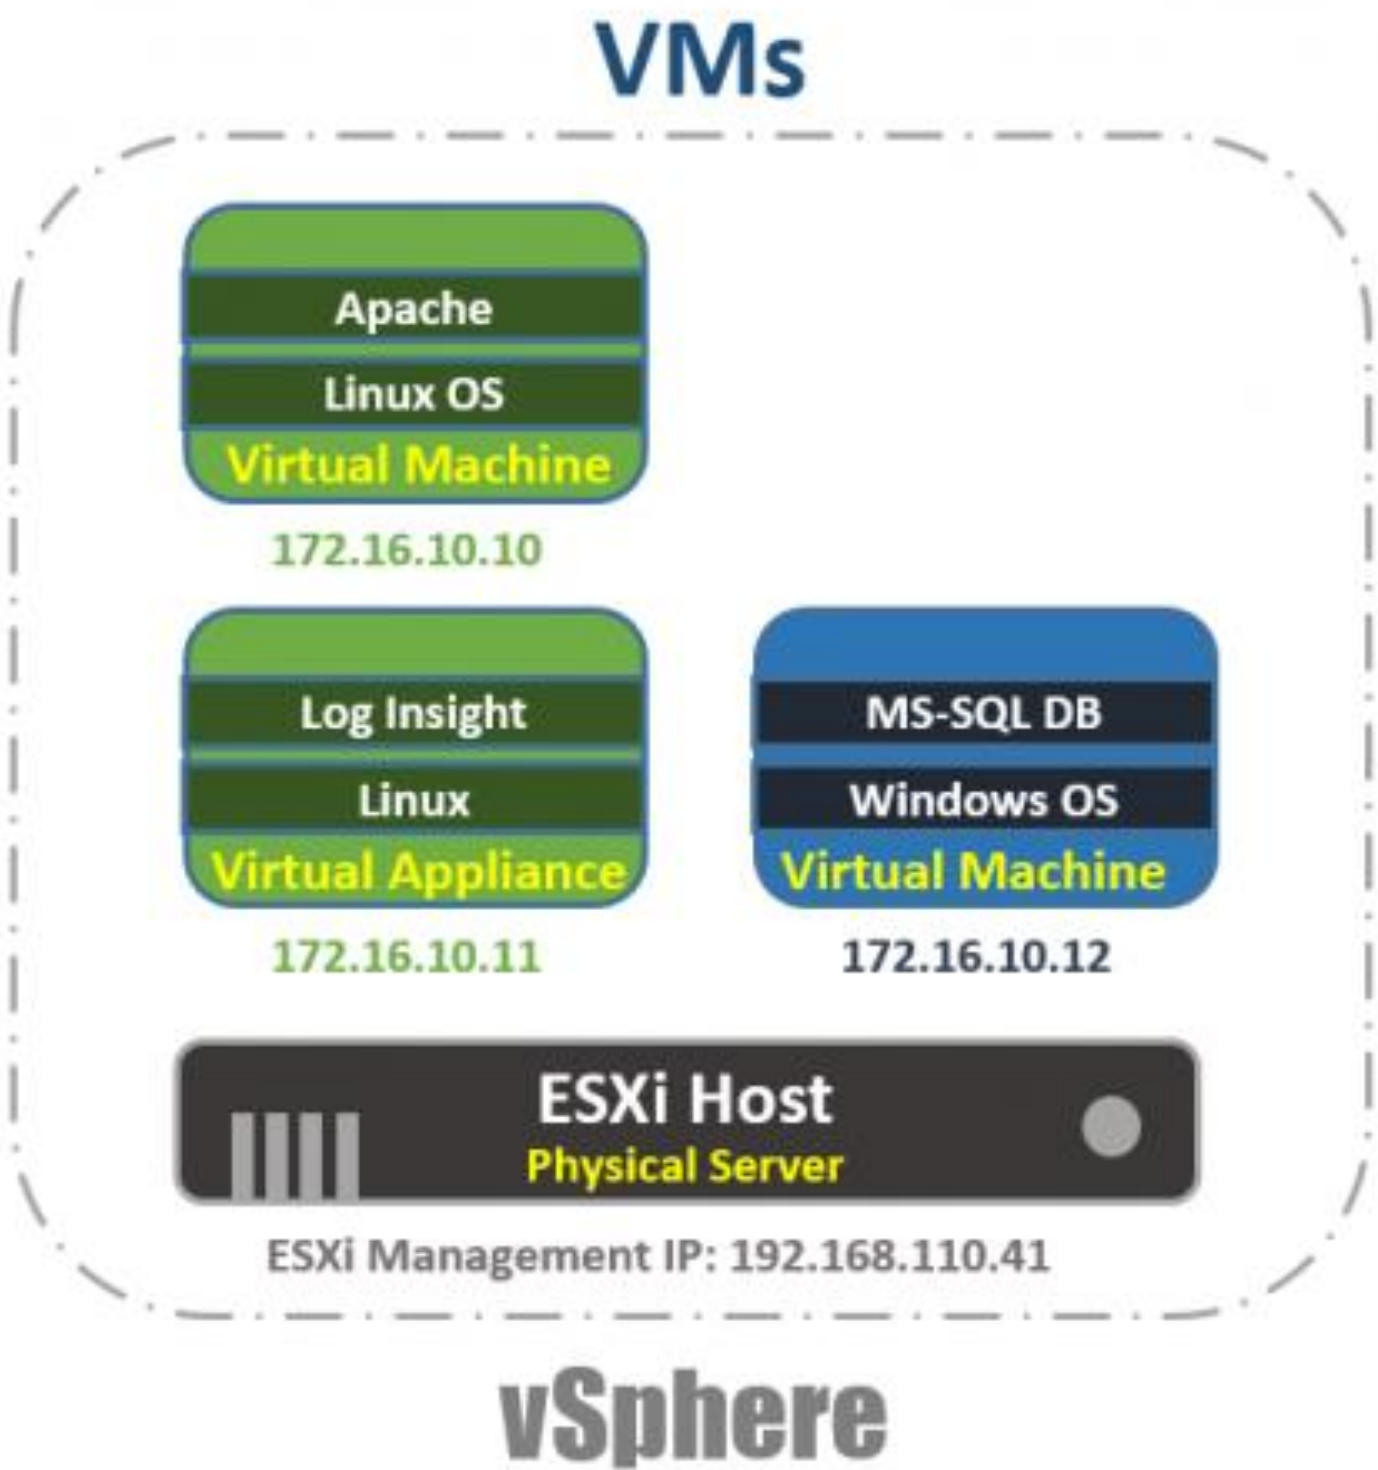
kubernetes

- It all started with Google growing and experiencing problems on managing the new scale of hardware and software
- The Datacentre as a Computer
(<https://research.google/pubs/pub35290/>)
- Abstract completely hardware (software defined datacentre)
- Abstract completely from network (software define network)
- Declarative application deployment (deploy is documentation)
- Self-Healing system based on desired state
- Ability to configure rules for automatic scaling
- Designed for multi-tenant
- Designed for integration (“API first” approach)

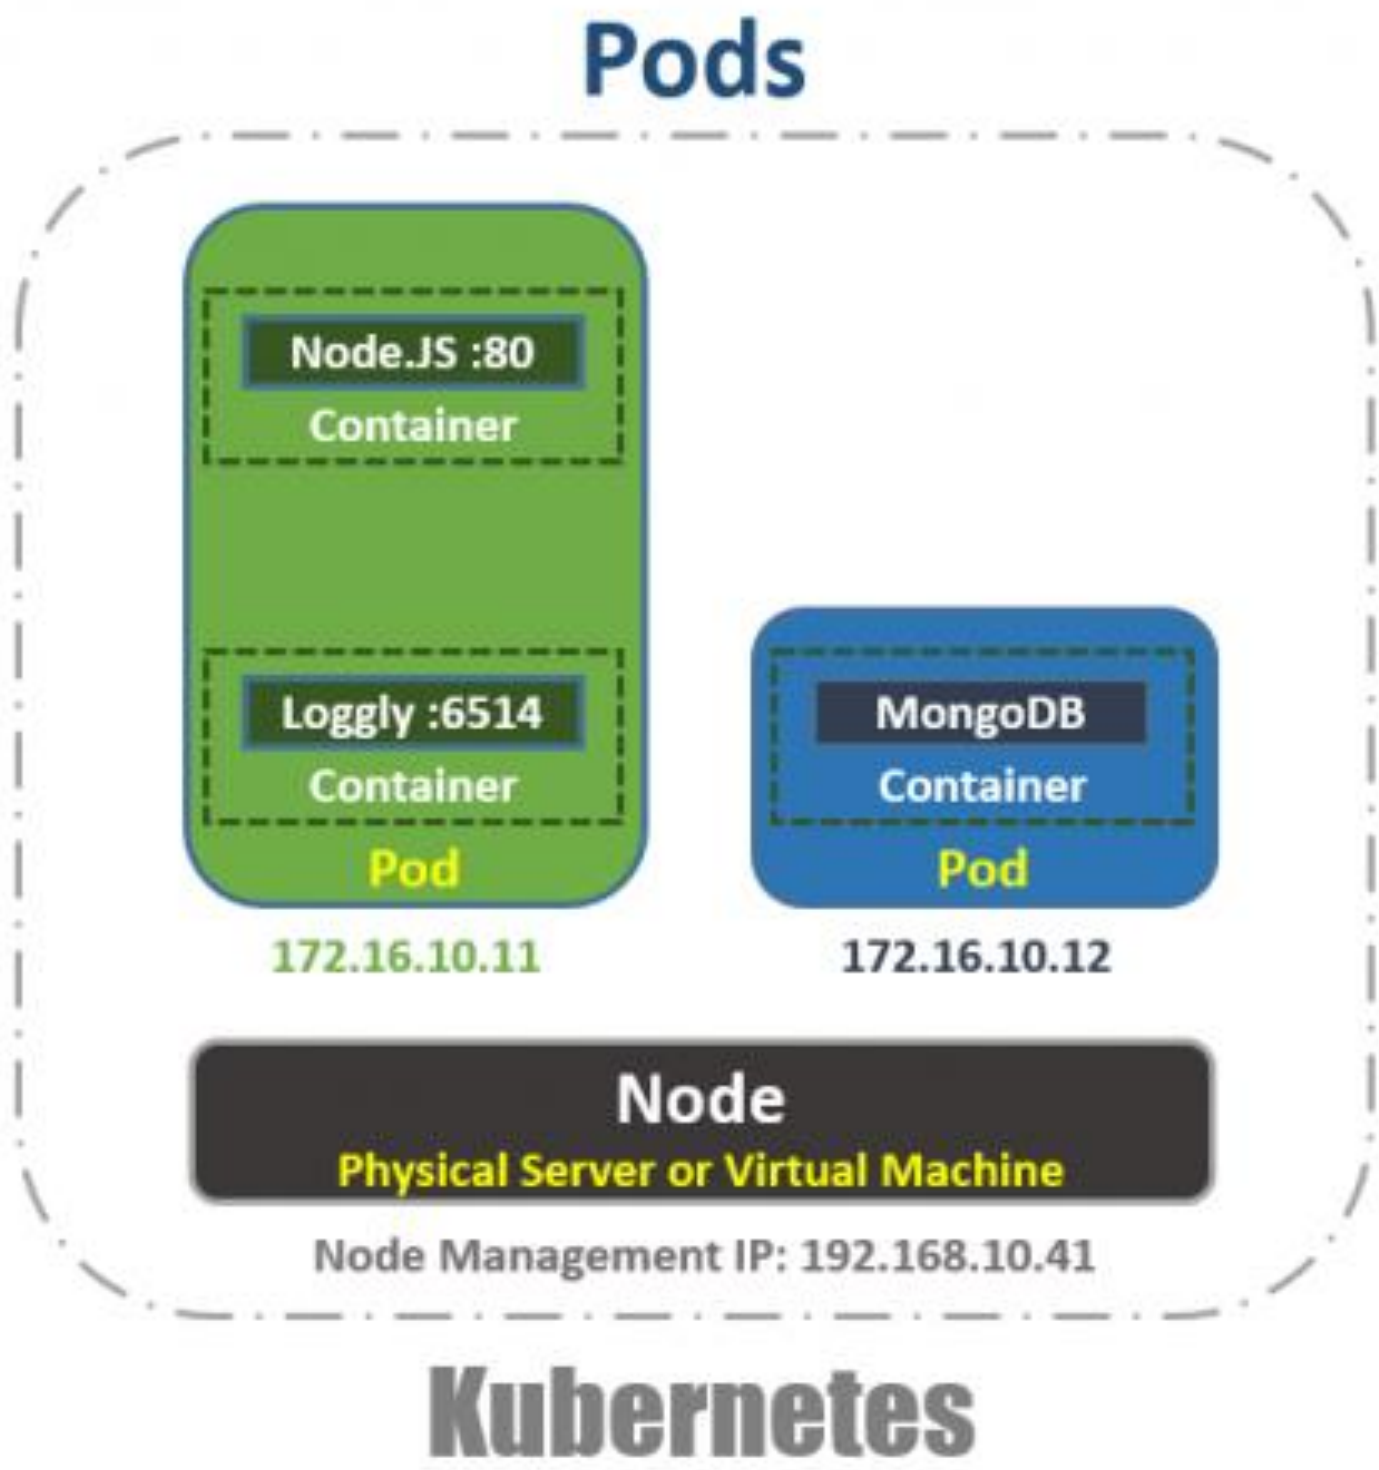
KUBERNETES VS VIRTUAL MACHINES



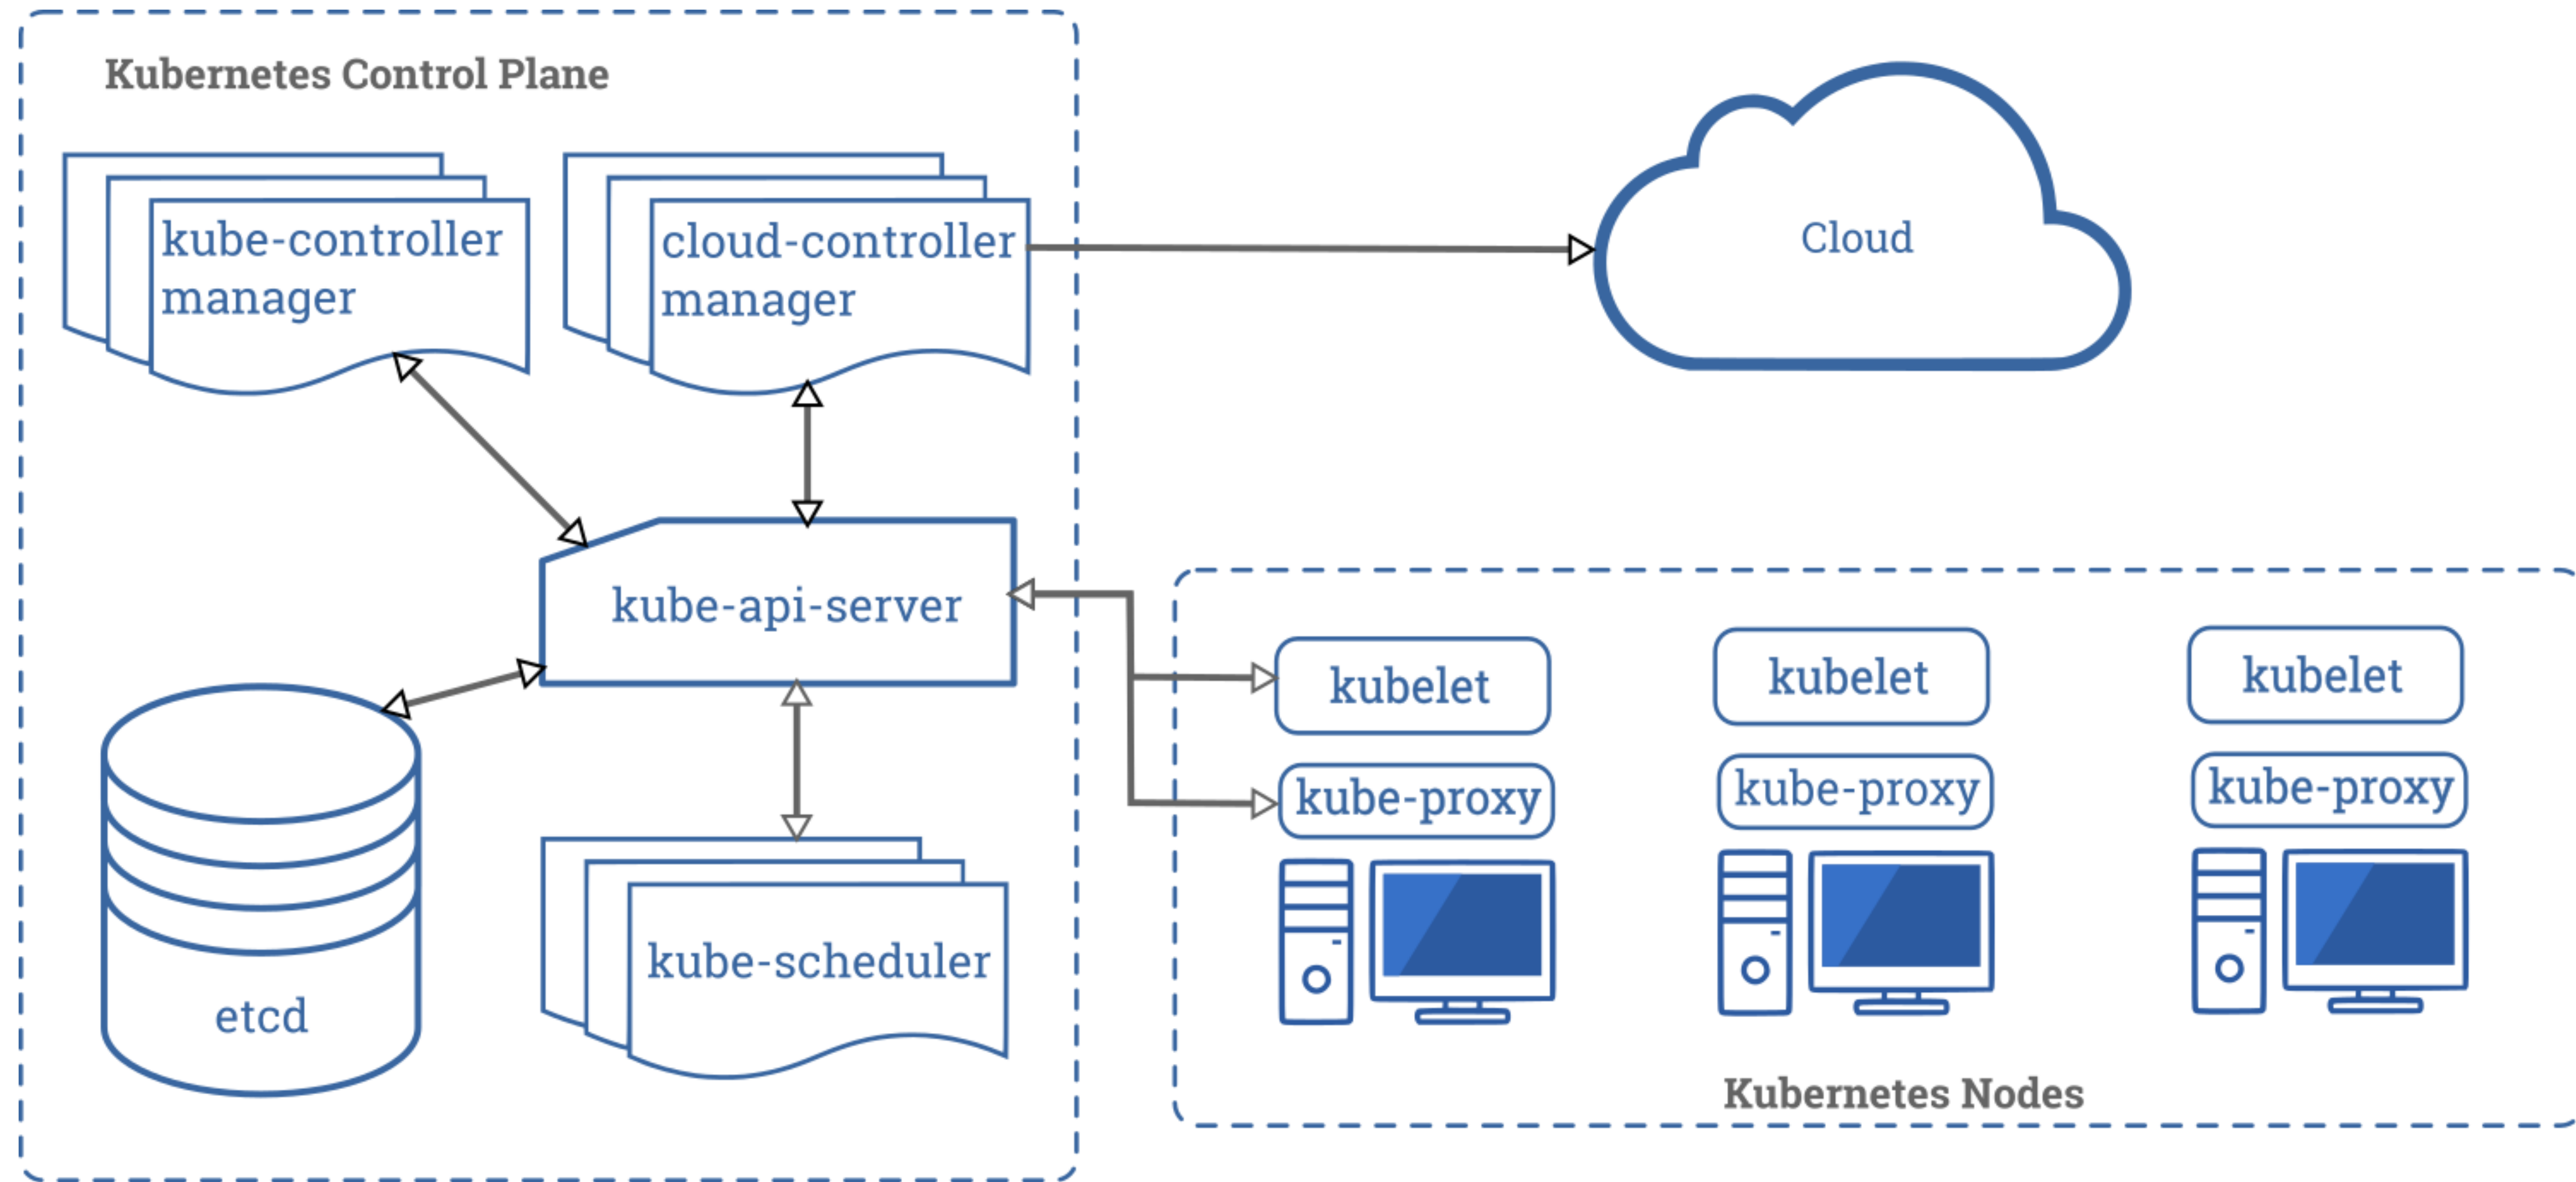
KUBERNETES VS VIRTUAL MACHINES



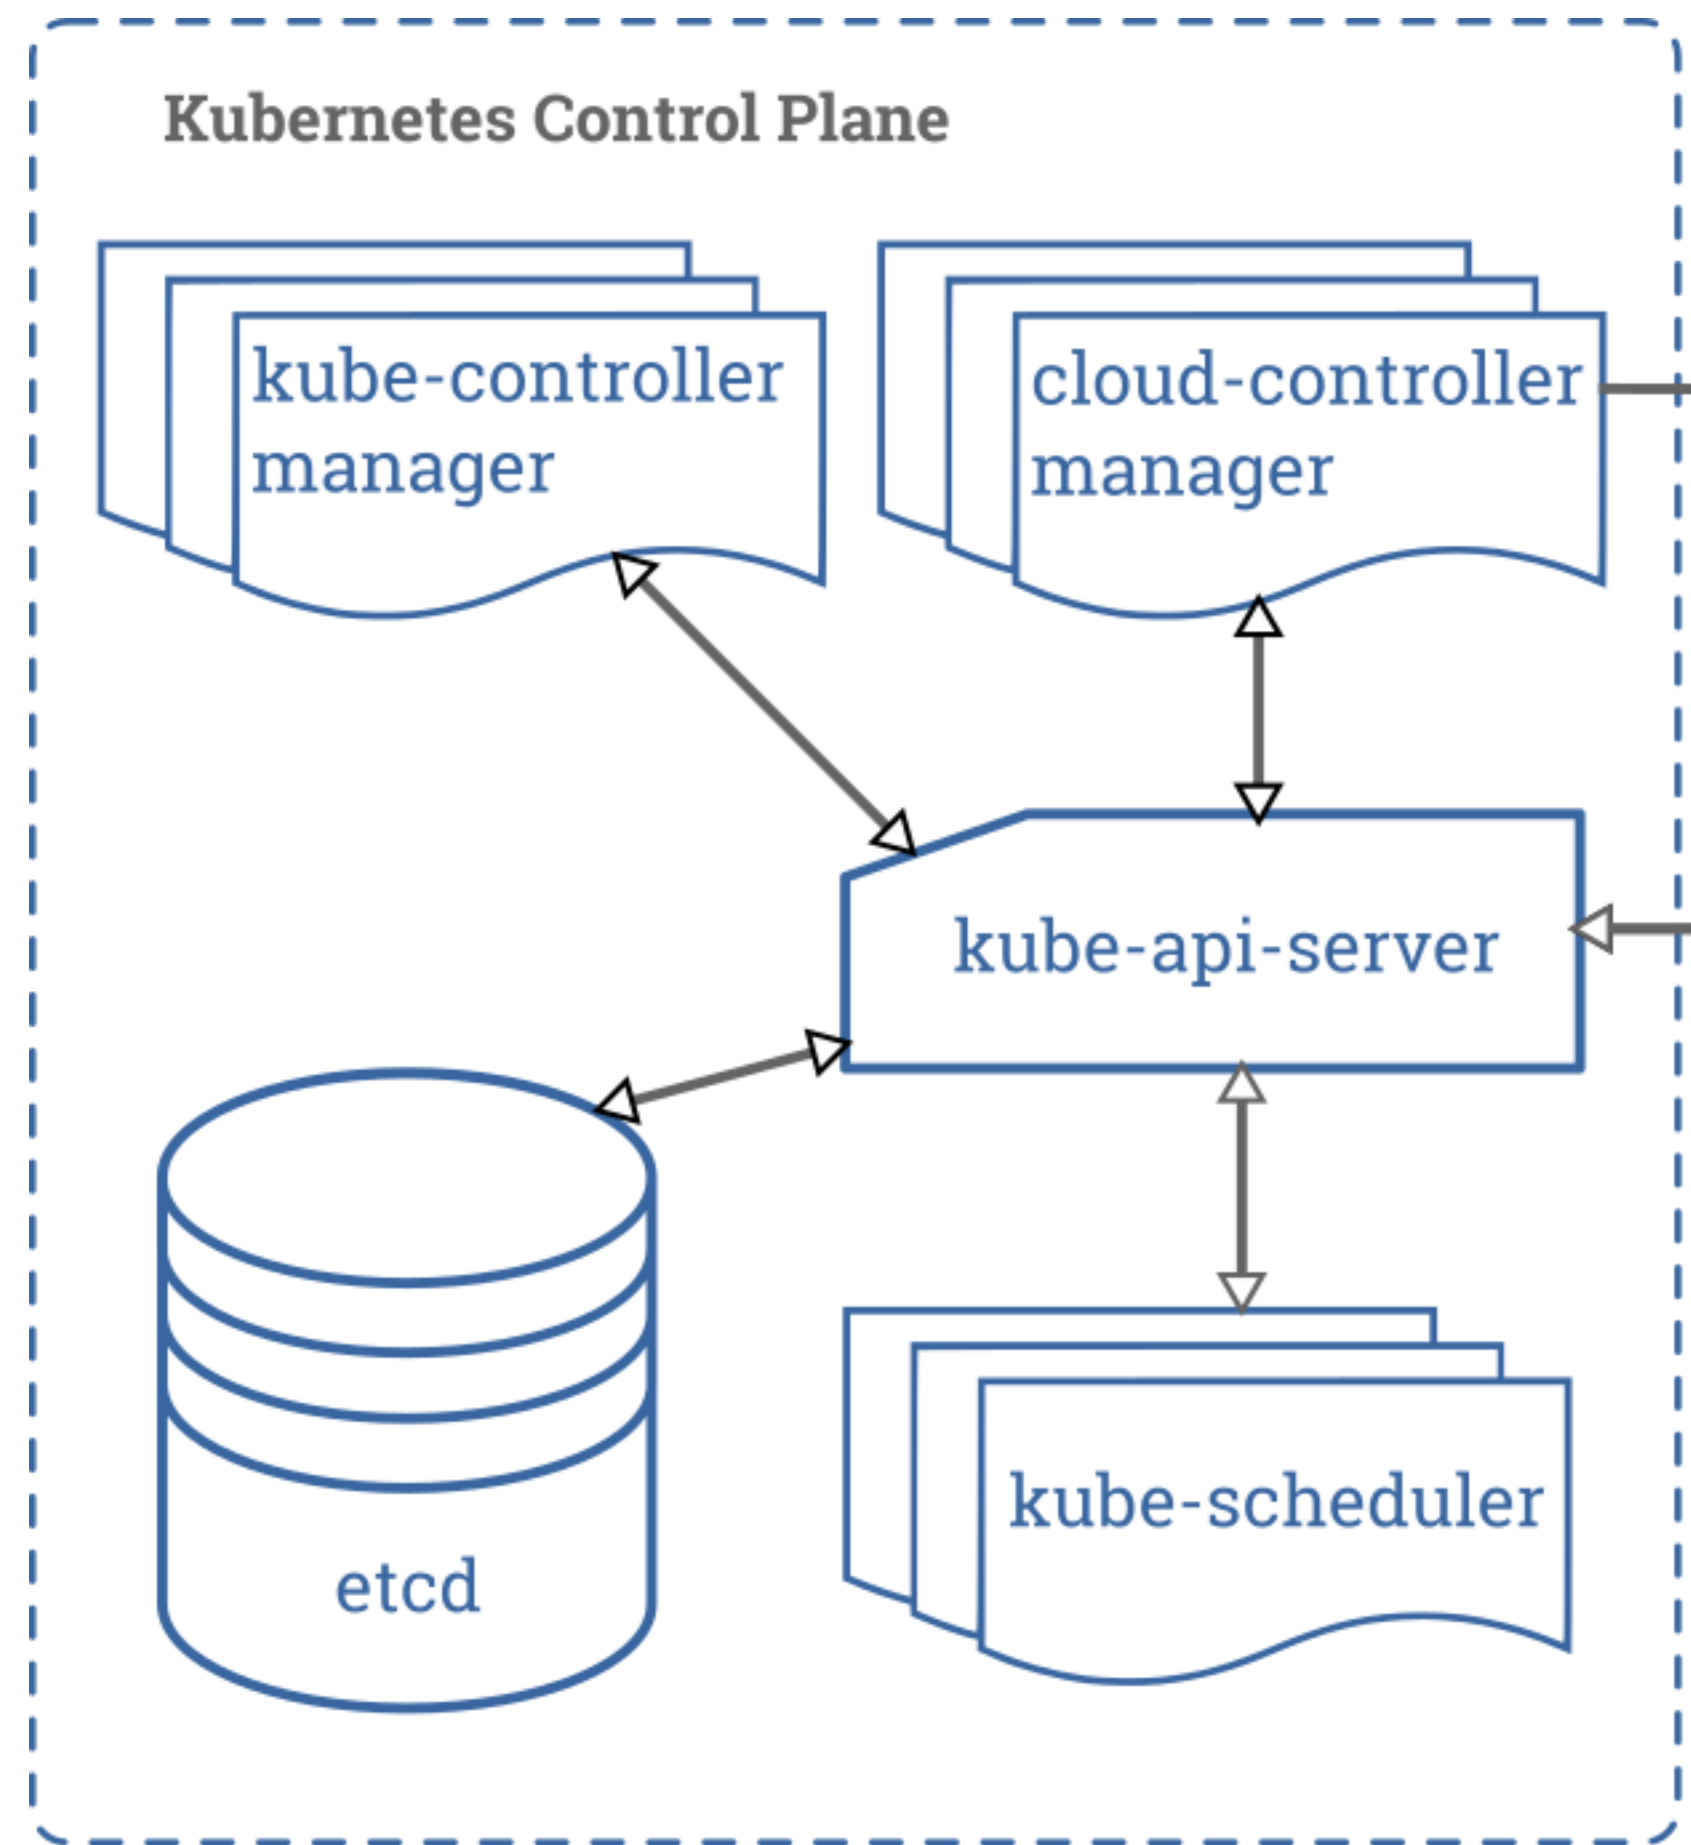
VS.



KUBERNETES ARCHITECTURE



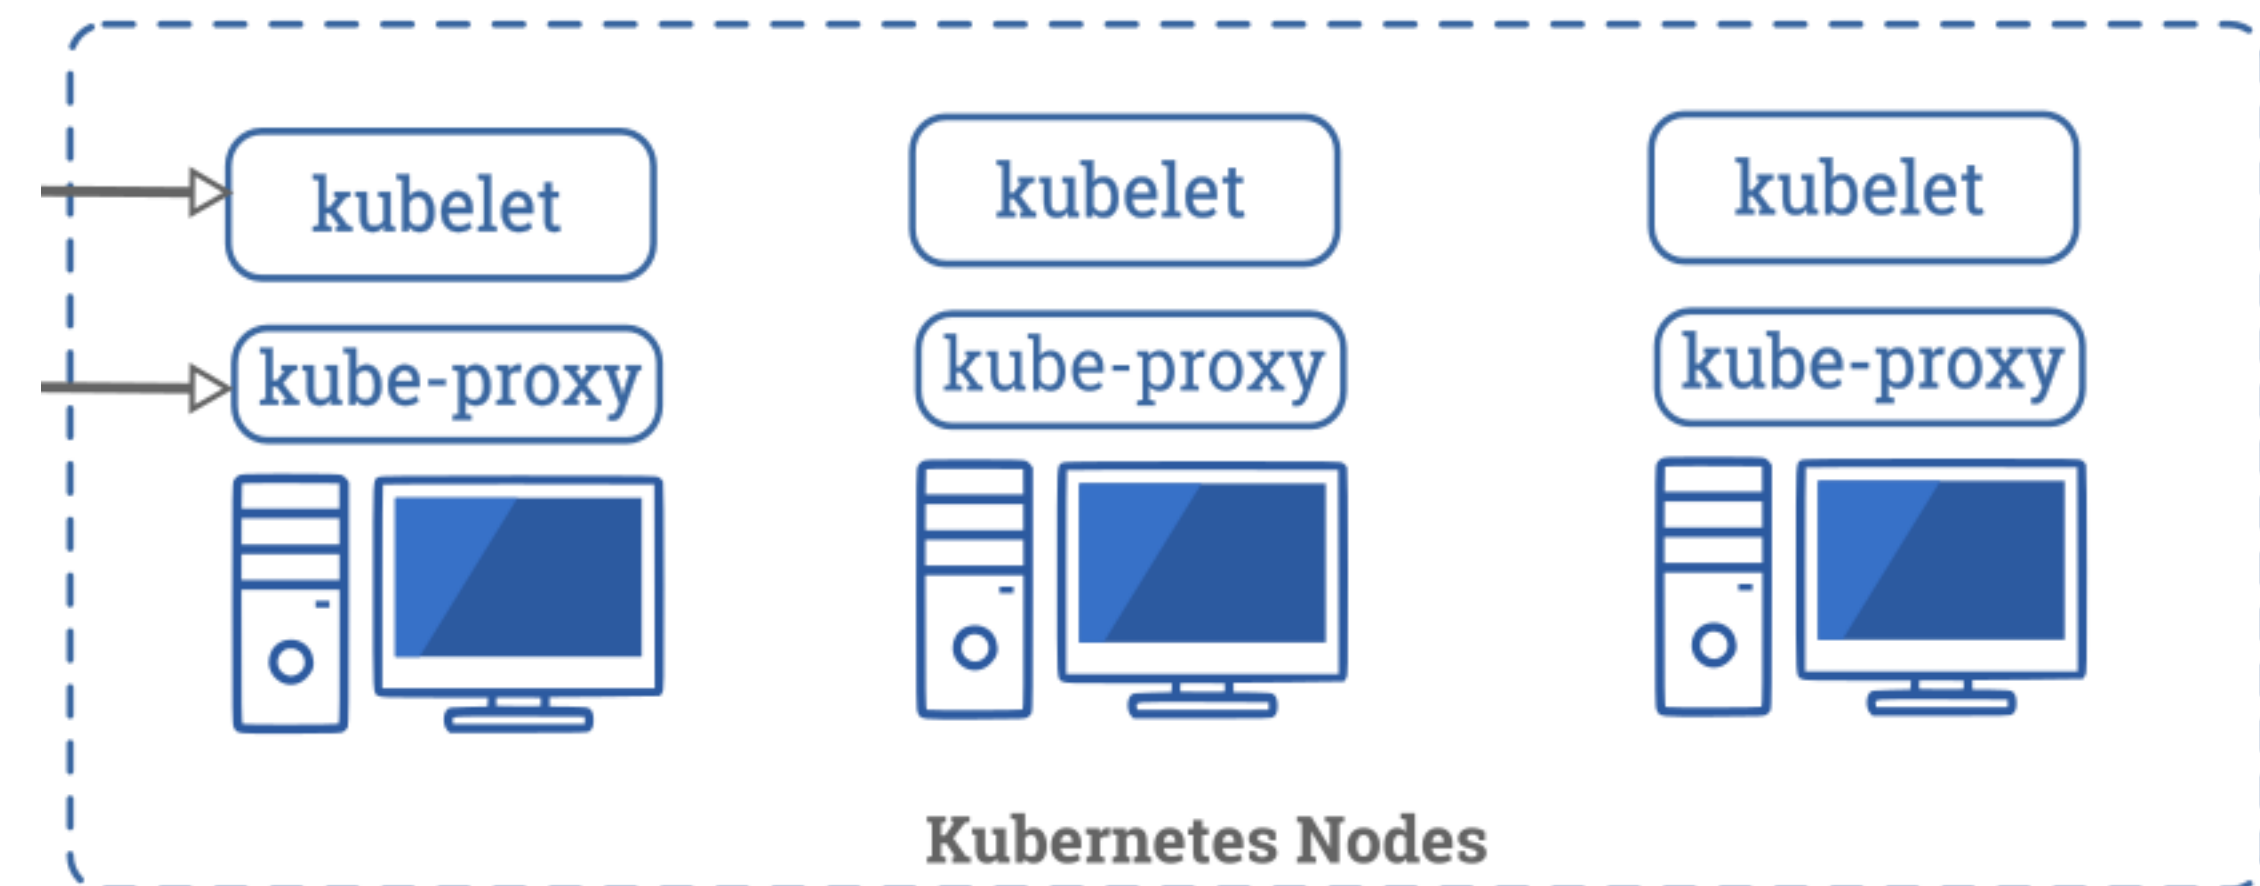
KUBERNETES ARCHITECTURE



- Etcd
 - The etcd project, developed by the team at CoreOS, is a lightweight, distributed key-value store that can be configured to span across multiple nodes.
 - Kubernetes uses etcd to store configuration data that can be accessed by each of the nodes in the cluster.
- Kube-apiserver
 - This is the main management point of the entire cluster as it allows a user to configure Kubernetes' workloads and organizational units
 - The API server implements a RESTful interface
- Kube-controller-manager
 - It manages different controllers that regulate the state of the cluster, manage workload life cycles, and perform routine tasks.
 - When a change is seen, the controller reads the new information and implements the procedure that fulfills the desired state.
- Kube-scheduler
 - The process that actually assigns workloads to specific nodes
 - The scheduler is responsible for tracking available capacity on each host to make sure that workloads are not scheduled in excess of the available resources.

KUBERNETES ARCHITECTURE

- ▶ Container Runtime
 - ▶ Typically Docker
 - ▶ Rkt and runC supported
- ▶ Kubelet
 - ▶ The kubelet service communicates with the master components to authenticate to the cluster and receive commands and work
 - ▶ The kubelet process then assumes responsibility for maintaining the state of the work on the node server.
- ▶ Kube-Proxy
 - ▶ To manage individual host subnetting and make services available to other components



KUBERNETES – THE BASIC CONCEPTS



kubernetes

- ▶ Cluster - A collection of hosts that aggregate their available resources including cpu, ram, disk, and their devices into a usable pool.
- ▶ Master - The master(s) represent a collection of components that make up the control plane of Kubernetes. These components are responsible for all cluster decisions including both scheduling and responding to cluster events.
- ▶ Node - A single host, physical or virtual capable of running pods. A node is managed by the master(s), and at a minimum runs both kubelet and kube-proxy to be considered part of the cluster.
- ▶ Namespace - A logical cluster or environment. Primary method of dividing a cluster or scoping access.



KUBERNETES – THE BASIC CONCEPTS

- Pod - A pod is the smallest unit of work or management resource within Kubernetes. It is comprised of one or more containers that share their storage, network, and context (namespace, cgroups etc).
- Deployment - A declarative method of managing stateless Pods and ReplicaSets. Provides rollback functionality in addition to more granular update control mechanisms.
- Service - Services provide a method of exposing and consuming L4 Pod network accessible resources. They use label selectors to map groups of pods and ports to a cluster-unique virtual IP.
- Volume - Storage that is tied to the Pod Lifecycle, consumable by one or more containers within the pod.
- ConfigMap - Externalized data stored within kubernetes that can be referenced as a commandline argument, environment variable, or injected as a file into a volume mount. Ideal for separating containerized application from configuration.
- Secret - Functionally identical to ConfigMaps, but stored encoded as base64, and encrypted at rest (if configured).

KUBERNETES — YAML FILES

! mlb-dev-public-site-graphql.yaml ✕

```
1  apiVersion: apps/v1beta1
2  kind: StatefulSet
3  metadata:
4    name: graphql-dapi
5  spec:
6    selector:
7      matchLabels:
8        app: graphql-dapi
9    serviceName: "graphql-dapi-svc"
10   replicas: 1
11   template:
12     metadata:
13       labels:
14         app: graphql-dapi
15     spec:
16       containers:
17         - name: graphql-dapi
18           image: mlbdevregistry.azurecr.io/mlbam-graphql-dapi:244580
19           imagePullPolicy: IfNotPresent
20           env:
21             - name: CULTURE_LIST
22               value: "en-us"
23             - name: DAPI_URL
24               value: "https://dapi.cms-dev.mlbinfra.com/"
25             - name: DAPI_VERSION
26               value: "v2"
27             - name : CUSTOM_ENTITIES
28               value: "players,videos,links,raffles,events,highlights,socialposts,promos,authors,lineups,playersext,hero,button,sidekick"
29           ports:
30             - containerPort: 4000
31       imagePullSecrets:
32         - name: mlbdevregistry.azurecr.io
33   ---
34   apiVersion: v1
35   kind: Service
36   metadata:
37     name: graphql-dapi-svc
38   spec:
39     type: LoadBalancer
40     ports:
41       - port: 80
42         targetPort: 4000
43     selector:
44       app: graphql-dapi
```

! mlb-dev-public-site-graphql.yaml

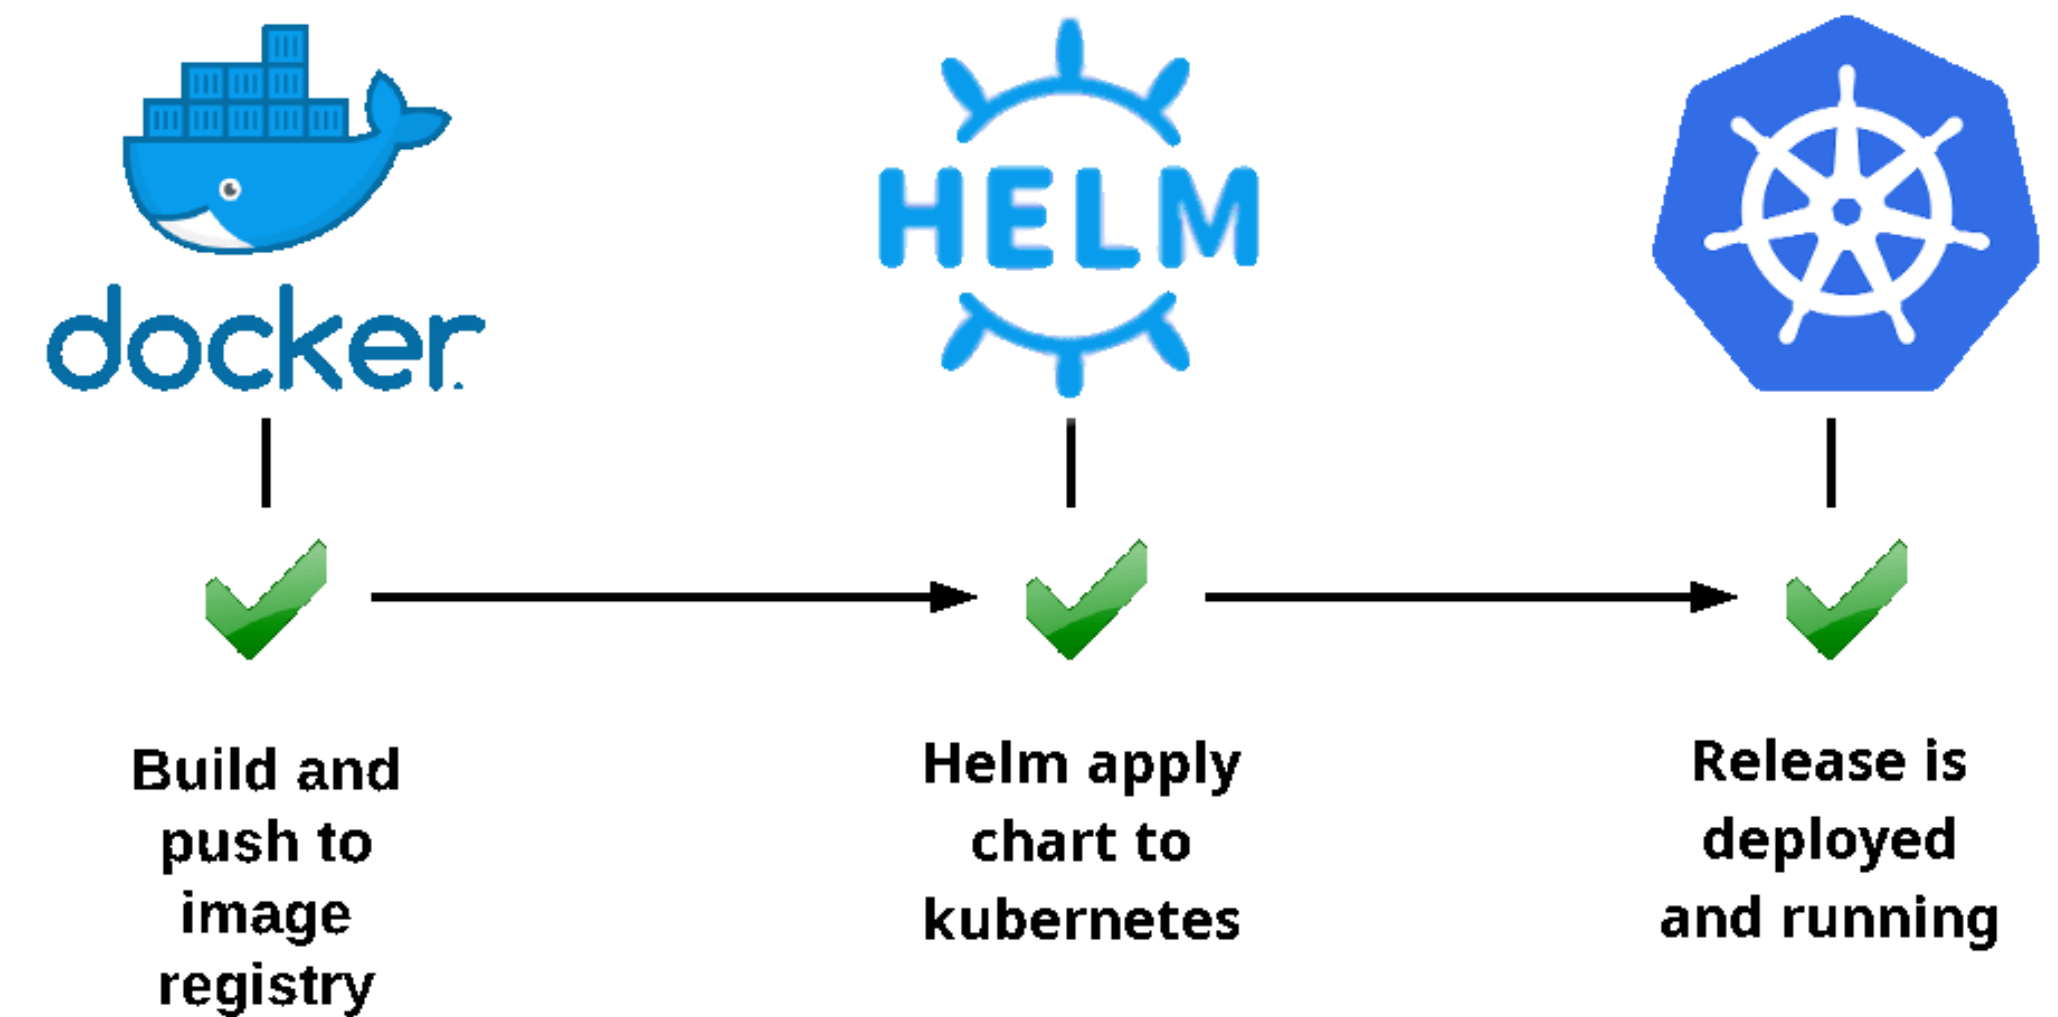
! mlb-dev-public-site-nginxproxy.yaml ✕

```
1  apiVersion: v1
2  kind: ConfigMap
3  metadata:
4    name: forge-proxy-nginx-conf
5  data:
6    nginx.conf: |
7      user nginx;
8      worker_processes 3;
9      error_log /var/log/nginx/error.log notice;
10     events {
11       worker_connections 10240;
12     }
13     http {
14       rewrite_log on;
15       log_format main
16         '$remote_addr:$remote_addr\t'
17         '$time_local:$time_local\t'
18         '$method:$request_method\t'
19         '$uri:$request_uri\t'
20         '$host:$host\t'
21         '$status:$status\t'
22         '$bytes_sent:$body_bytes_sent\t'
23         '$referer:$http_referer\t'
24         '$useragent:$http_user_agent\t'
25         '$forwardedfor:$http_x_forwarded_for\t'
26         '$request_time:$request_time';
27       access_log /var/log/nginx/access.log main;
28
29       server {
30         listen 80;
31
32         location ~ ^/cms/api/(filters|menuitems|views) {
33           rewrite ^/cms/api/(.*)$ /$1 break;
34           proxy_pass https://siteassets-api.cms-dev.mlbinfra.com;
35           proxy_pass_request_headers on;
36           proxy_redirect off;
37         }
38         location ~ ^/cms/api/(layouts|modules|templates) {
39           proxy_pass http://qa.mlbstatic.com;
40           proxy_pass_request_headers on;
41           proxy_redirect off;
42         }
43         location ~ ^/cms/api/(routes|pages|public|working)/ {
44           rewrite ^/cms/api/(.*)$ /$1 break;
45           proxy_pass https://routing-api.cms-dev.mlbinfra.com;
46           proxy_pass_request_headers on;
47           proxy_redirect off;
48         }
49         location ~ ^/cms/api/(echo|resources|data) {
50           rewrite ^/cms/api/(.*)$ /$1 break;
```


and more

HELM PACKAGE MANAGER

- Part of the Cloud Native Computing Foundation
- Designed to simplify management of dependencies on Kubernetes deployments
- CHARTS: Helm packages, a few YAML configurations files
- Mostly standard Kubernetes YAML format
- Templates and Values yaml files used to abstract composition of Kubernetes YAML files with variables (e.g. by environment)
- Requirements.yaml used to define dependencies



HELM CHART EXAMPLE

ott-entitlement

templates

blue

ott-entitlement-config.yaml

ott-entitlement-deployment.yaml

ott-entitlement-hpa.yaml

green

ott-entitlement-backoffice-config.yaml

ott-entitlement-backoffice-deployment.yaml

ott-entitlement-backoffice-ingress.yaml

ott-entitlement-backoffice-service.yaml

ott-entitlement-ingress.yaml

ott-entitlement-secrets.yaml

ott-entitlement-service-loadbalancer.yaml

ott-entitlement-service.yaml

ott-entitlement-viprules-secret.yaml

Chart.yaml

values.yaml

ott-entitlements-provider

ott-forge-catalog-connector

ott-identity-adapter-adobe

ott-identity-management

ott-ingenico-adapter

You updated terraform 4 hours ago — Create a pull request

ContentsHistoryCompareBlameEditRenameDeleteDownload

1

2

3

4

5

6

7

8

9

10

11

12

13

14

15

16

17

18

19

20

21

22

23

24

25

26

27

28

29

30

31

32

33

34

35

36

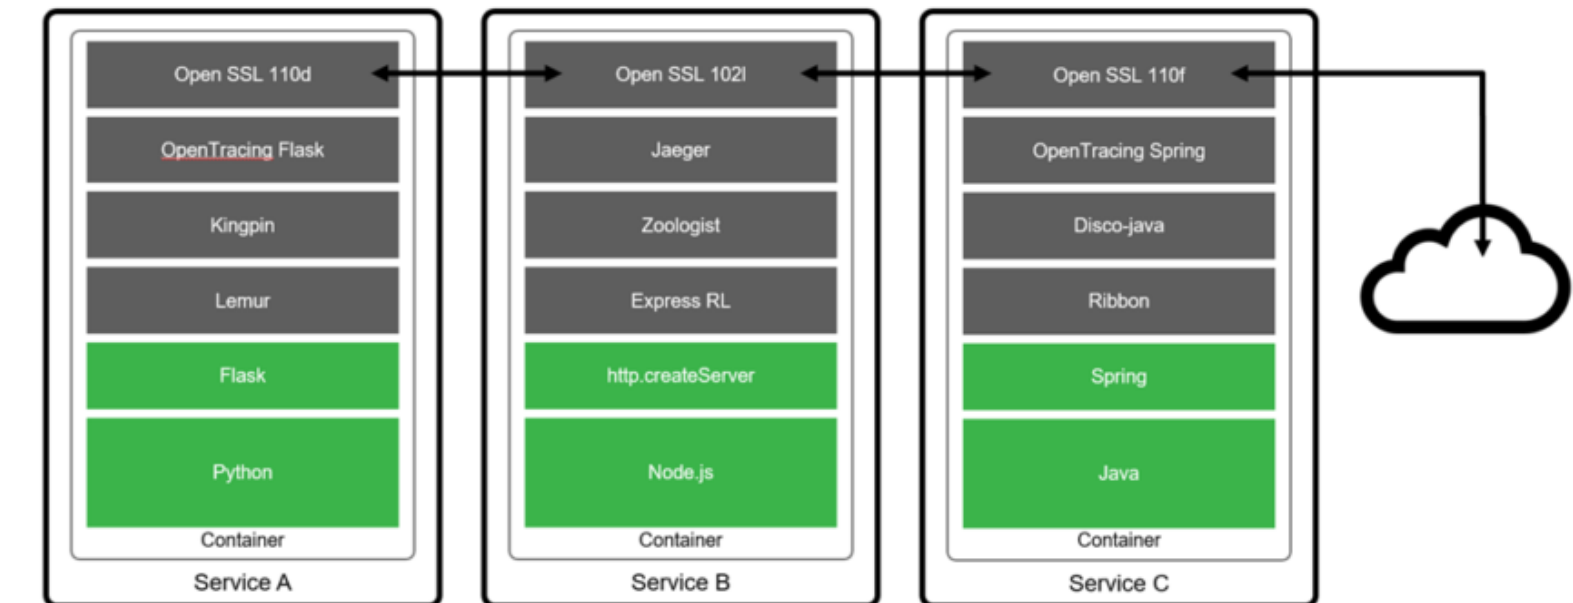
```
1 {{ if .Values.component.backoffice.enabled }}
2 apiVersion: extensions/v1beta1
3 kind: Deployment
4 metadata:
5   name: {{ .Chart.Name }}-bo
6   namespace: {{ .Values.amplify.namespace }}
7   labels:
8     name: {{ .Chart.Name }}-bo
9 spec:
10   replicas: 1
11   strategy:
12     type: RollingUpdate
13     rollingUpdate:
14       maxSurge: 1
15   template:
16     metadata:
17       namespace: {{ .Values.amplify.namespace }}
18       labels:
19         name: {{ .Chart.Name }}-bo
20         tier: dotnet-service
21   spec:
22     containers:
23     - name: {{ .Chart.Name }}-bo
24       {{ if regexMatch "[a-zA-Z]+" .Chart.AppVersion }}
25       image: ottregistrydev.azurecr.io/ott-entitlement-bo:{{ .Chart.AppVersion }}
26       {{ else }}
27       image: ottregistryprd.azurecr.io/ott-entitlement-bo:{{ .Chart.AppVersion }}
28       {{ end }}
29     resources:
30       requests:
31         memory: "{{ .Values.component.backoffice.resources.requests.memory }}"
32         cpu: "{{ .Values.component.backoffice.resources.requests.cpu }}"
33       limits:
34         memory: "{{ .Values.component.backoffice.resources.limits.memory }}"
35         cpu: "{{ .Values.component.backoffice.resources.limits.cpu }}"
36   ports:
```



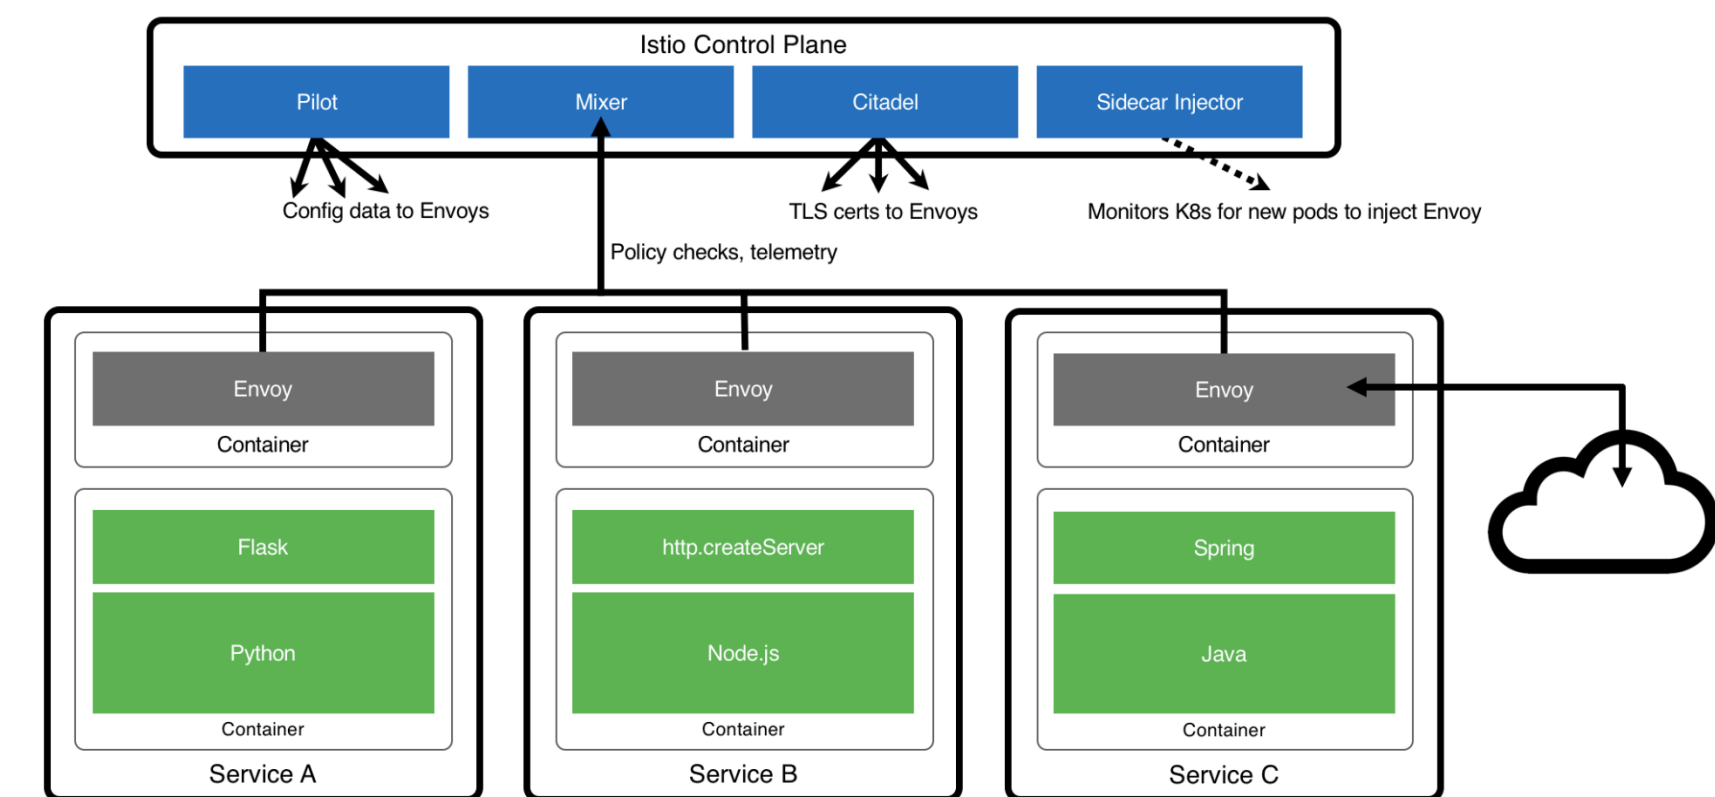
ISTIO – SERVICE MESH

- Traffic Management
 - Decouples traffic flow and infrastructure scaling, letting you specify via Pilot what rules you want traffic to follow rather than which specific pods/VMs
- Security
 - Strong identity, powerful policy, transparent TLS encryption, and authentication, authorization and audit (AAA) tools
- Policy and Telemetry
 - A flexible model to enforce authorization policies and collect telemetry for the services in a mesh
- Performance and Scalability
 - Support for Horizontal Pod Autoscaling

Managing Microservices Without a Service Mesh



Managing Microservices With Istio





LINKS

Processes, Containers, Virtual Machines -
<https://medium.com/@jessgreb01/what-is-the-difference-between-a-process-a-container-and-a-vm-f36ba0f8a8f7>

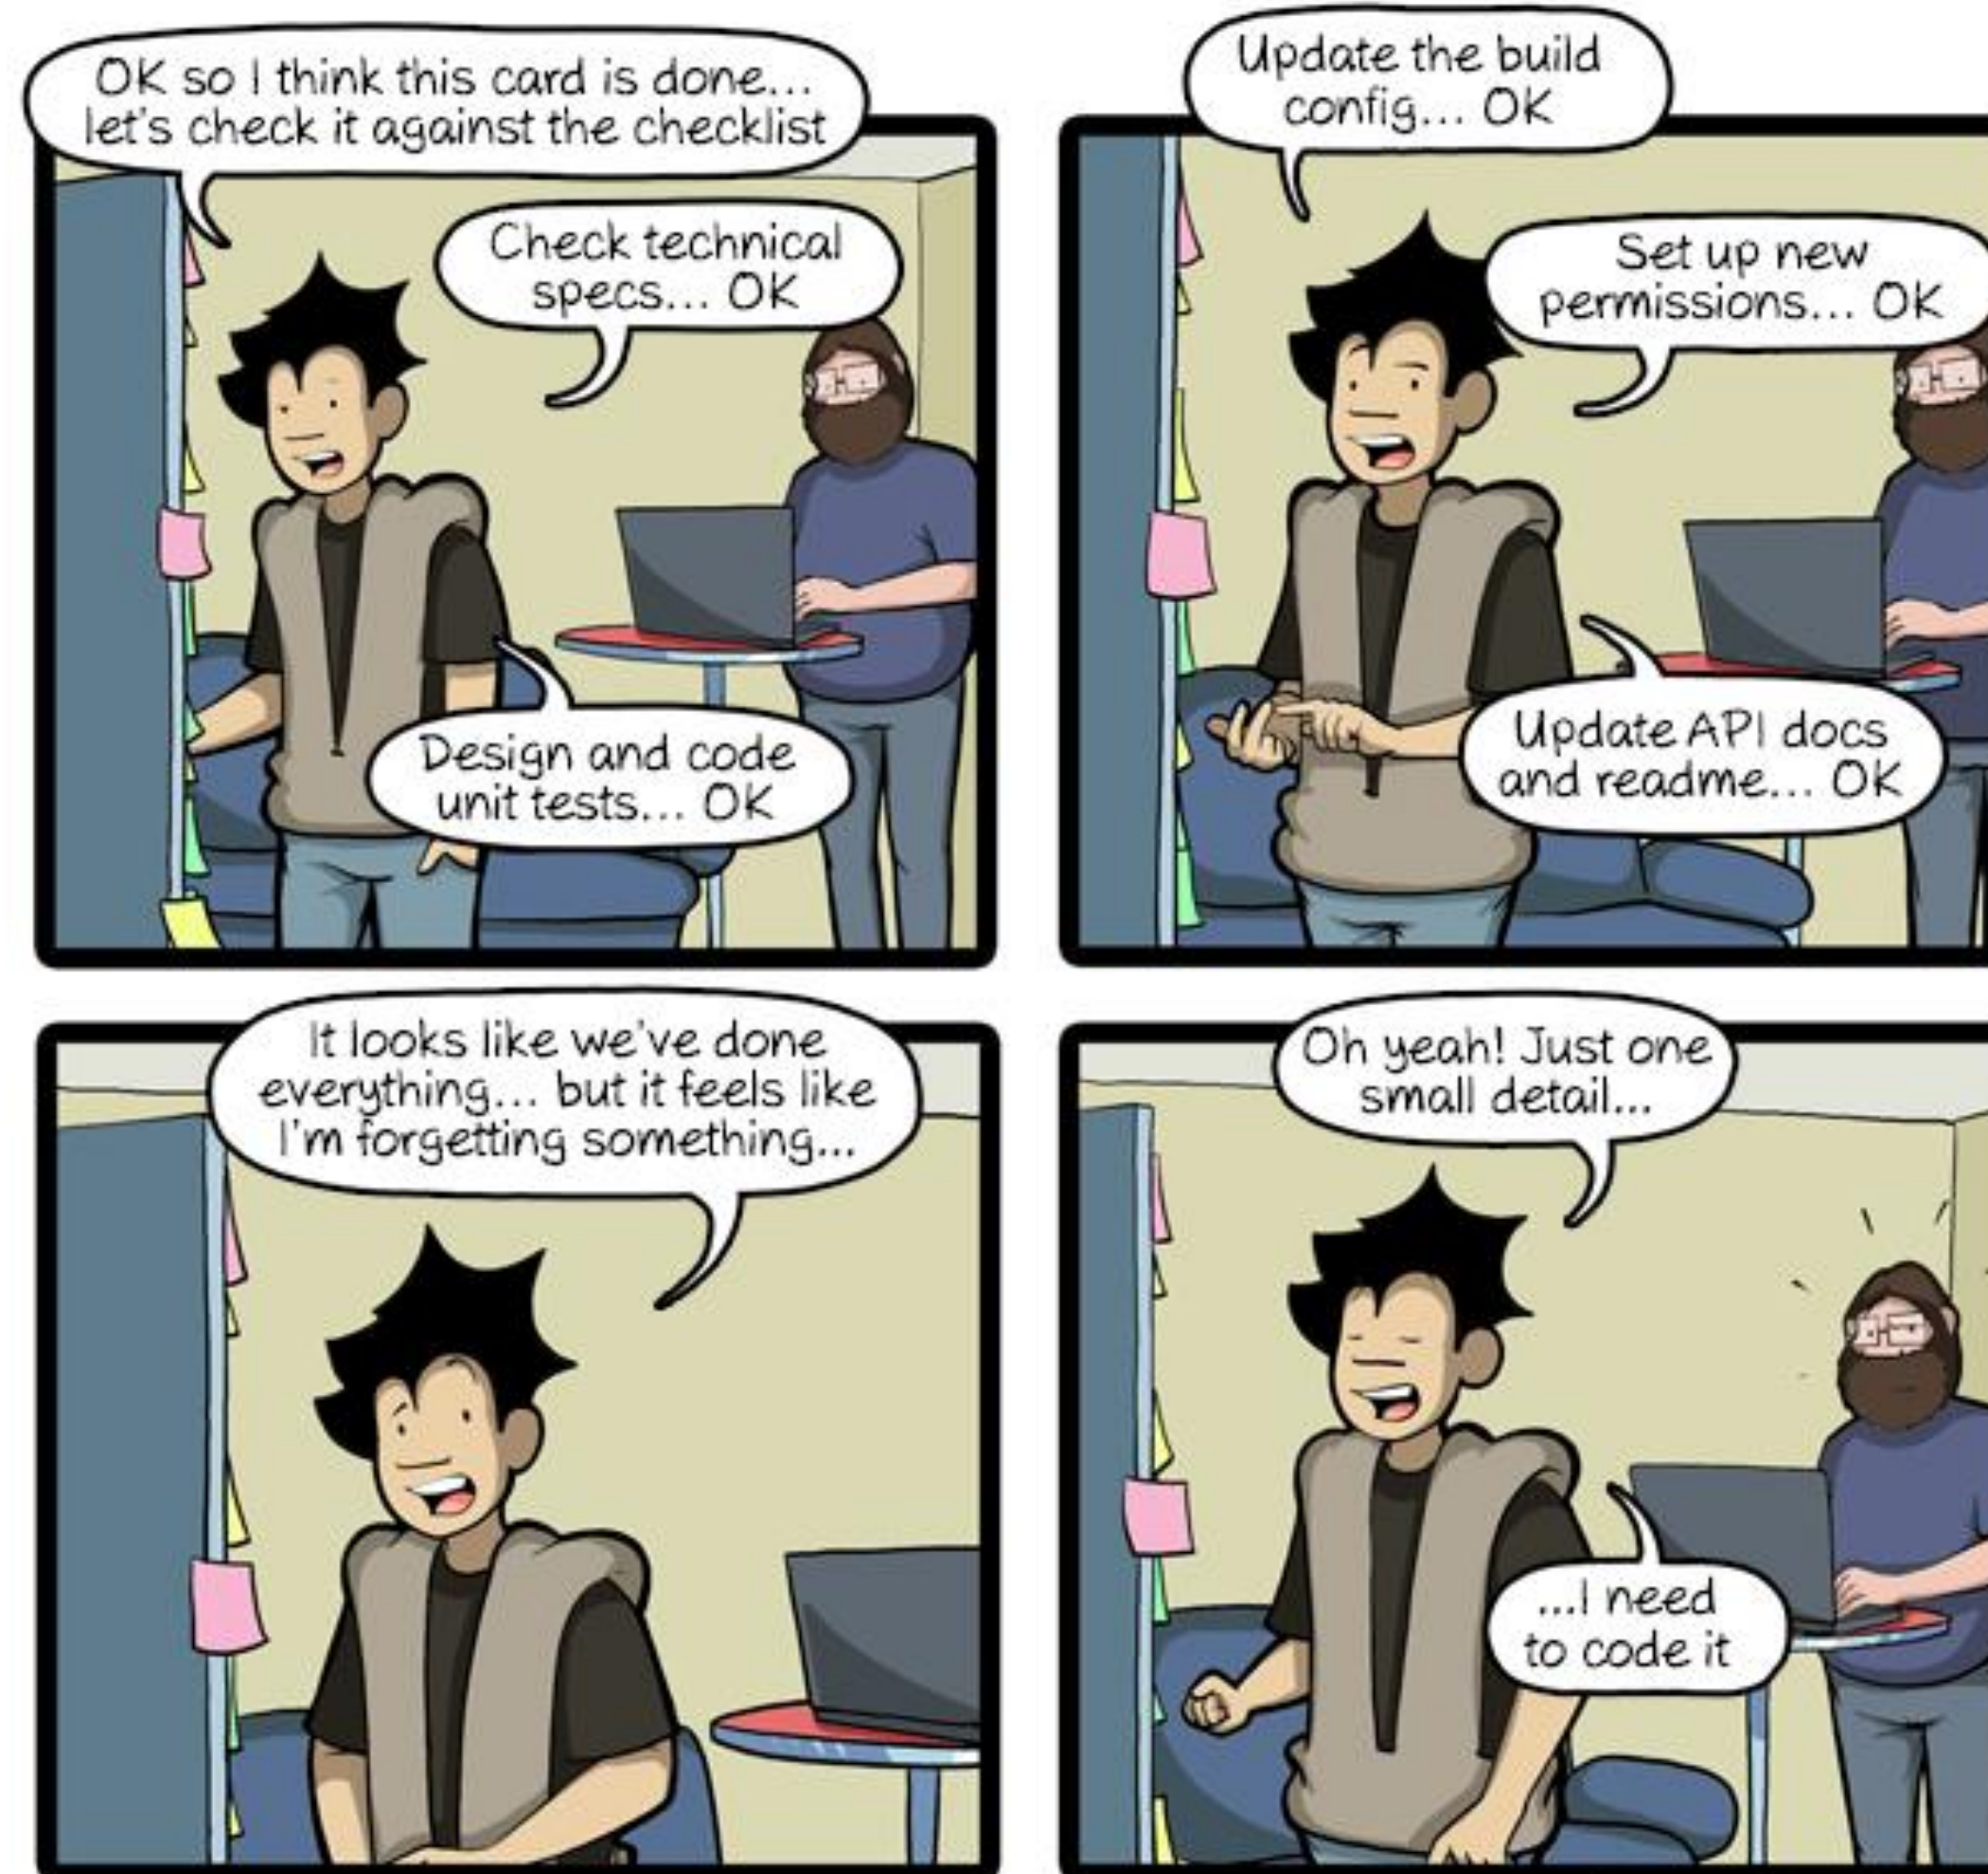
Introduction to Kubernetes for Vmware users -
<https://blogs.vmware.com/cloudnative/2017/10/25/kubernetes-introduction-vmware-users/>

Introduction to Kubernetes Architecture - <https://phoenixnap.com/kb/understanding-kubernetes-architecture-diagrams>

Docker and Windows -
<https://techcommunity.microsoft.com/t5/windows-dev-appconsult/first-steps-with-docker-introduction/ba-p/317547>

Kubernetes and Windows -
<https://techcommunity.microsoft.com/t5/windows-dev-appconsult/first-steps-with-docker-and-kubernetes-introduction/ba-p/357525>

THE END – Q&A ?



CommitStrip.com